

CONVOLUTIONAL NEURAL NETWORKS BASED PNEUMONIA DIAGNOSIS

*A Graduate Project Report submitted to Manipal Academy of Higher
Education in partial fulfilment of the requirement for the award of the
degree of*

BACHELOR OF TECHNOLOGY

In

Electronics and Communication Engineering

Submitted by

Shreya K. R

Reg. No: 140907759

Under the guidance of

Mr Vijay Kumar Gupta
Deputy General Manager
Samsung Research Institute, Noida

Ms Navya K T
& Assistant Professor, Senior Scale
Dept. of Electronics and Comm.



DEPARTMENT OF ELECTRONICS AND COMMUNICATION ENGINEERING

MANIPAL INSTITUTE OF TECHNOLOGY

(A Constituent Institution of Manipal Academy of Higher Education)

MANIPAL – 576104, KARNATAKA, INDIA

JULY 2018



DEPARTMENT OF ELECTRONICS AND COMMUNICATION ENGINEERING

MANIPAL INSTITUTE OF TECHNOLOGY

(A Constituent Institution of Manipal Academy of Higher Education)

MANIPAL – 576 104 (KARNATAKA), INDIA

Manipal
12.07.2018

CERTIFICATE

This is to certify that the project titled **CONVOLUTIONAL NEURAL NETWORKS BASED PNEUMONIA DIAGNOSIS** is a record of the bonafide work done by **SHREYA K.R** (Reg. No. 140907759) submitted in partial fulfilment of the requirements for the award of the Degree of Bachelor of Technology (BTech) in **ELECTRONICS AND COMMUNICATION ENGINEERING** of Manipal Institute of Technology, Manipal, Karnataka, (A Constituent Institution of Manipal Academy of Higher Education), during the academic year 2017 - 2018.

Ms. Navya K T

*Assistant Professor, Senior Scale
E & C, M.I.T MANIPAL*

Prof. Dr. M. Sathish Kumar

*HOD, E & C.
M.I.T, MANIPAL*

ACKNOWLEDGMENTS

I would like to thank Dr. D Srikanth Rao, Director, MIT Manipal for providing a strong academic environment, with an appreciation for innovation, that allowed me to take on this project. I would like to extend my gratitude to Dr. M Sathish Kumar, Head of Department, Department of Electronics and Communication, MIT Manipal for the unparalleled support, facilities and infrastructure the department provides.

I would like to express my sincere gratitude towards Mr Vijay Kumar Gupta, Deputy General Manager, Samsung Research Institute, Noida for his guidance, supervision and above all his approachability throughout the duration of this project. I am also thankful to Mr Ujjawal Rastogi, Chief Engineer, Project Leadership team for the continuous support, feedback and constructive criticism he provided while the project was being developed.

I am extremely grateful to my project guide, Ms Navya K T, Assistant Professor- Senior Scale, MIT Manipal, for her patience and insightful suggestions that allowed me to build this project. I thank her for her support and time.

ABSTRACT

Data is continuously generated at a fast pace, from various sources. The sheer volume of this constantly generated data leads it to be called big data. Big data in itself might be meaningless, but when processed, can hold very important trends and patterns. A large amount of this big data is in the form of images, which are particularly hard to process due to their high dimensionality and matrix representation which robs the image of its true meaning. Processing image type big data using artificial intelligence algorithms has caught on very rapidly. Medical Diagnostics is an area where there are a large number of scans available, which can be put to technical use. This project aims to incorporate artificial intelligence into Samsung digital radiography devices to allow the device to diagnose whether a lung X-Ray belongs to a healthy patient, or one infected with Pneumonia.

This is done by training a convolutional neural network using a labelled open source dataset which has 5216 lung X-Rays of both healthy and Pneumonia infected patients. The network is modified by tuning the hyper-parameters and varying the network architecture till a satisfactory accuracy of prediction is achieved. The project starts by training CNNs with simpler datasets like MNIST handwritten characters and Hortensia flowers to understand the nuances of CNNs, gradually building up to the final larger architecture for diagnosing X-Rays with Pneumonia.

The network is trained and modified and trained again till it achieves a high validation accuracy of 90%. There is slight over-fitting of the network to the training data but this cannot be avoided due to the large number of parameters in the network to overcome the low interclass variation and high intra-class variation. The network is able to make predictions for the test images to an accuracy of 94%. The deviation from the correct predictions can be explained on the basis of visual cues. The accuracy is seen to be growing throughout the training to plateau towards the end, whereas the model loss decreases.

While the project achieves a good validation accuracy and can be deployed into digital radiography computational systems, the scope of the project can be increased with an increase in the size and variety of the datasets. This project uses an existing open dataset, however, there can be a link created with local hospitals to collect scans for research purposes, to create more comprehensive datasets. Incorporating artificial intelligence into diagnostics can be the next frontier in medicine. It can greatly reduce the cases of misdiagnosis as CNNs become more powerful with comprehensive datasets.

LIST OF TABLES

Table No	Table Title	Page No
1.1	Pneumonia Death Statistics in Children	4
1.2	Project Schedule	6

LIST OF FIGURES

Figure No	Figure Title	Page No
1.1	Image represented as a matrix	1
1.2	Skin cancer image classification	2
1.3	Object recognition by Facebook	3
1.4	Google super-enhancement	3
2.1	Structure of a computational neuron	8
2.2	Fully connected feed forward Neural Network	9
2.3	RGB splitting of an image	10
2.4	A convolutional neural network	10
2.5	A convolution operation on image	11
2.6	Zero padding an image	11
2.7	a) Input Image b) Kernel c) Feature map	12
2.8	Sigmoid and ReLu activation	13
2.9	Max pooling	13
2.10	Binary cross-entropy loss function	14
2.11	Gradient descent optimization	15
2.12	Dropout in a neural network	16
2.13	Distortion and occlusion	17
2.14	Intra-class Variation, Viewpoint Variation, Illumination, Background Clutter	18
2.15	Normal and Pneumonia Lung X-Ray	19
3.1	Flowchart of MNIST CNN Methodology	20
3.2	MNIST Sample	21
3.3	Test and training data sizes	21
3.4	Model Summary of MNIST CNN	22
3.5	Testing MNIST CNN	23
3.6	Flowchart of X-Ray CNN Methodology	24
3.7	Hortensia data size	25
3.8	X-Ray data size	25
3.9	Hortensia dataset samples	26
3.10	Hortensia sample sizes	26
3.11	X-Ray dataset samples	27
3.12	X-Ray sample sizes	27
3.13	Hortensia CNN model summary	28
3.14	X-Ray CNN model summary	29
3.15	Change of learning rate	30

Figure No	Figure Title	Page No
4.1	Accuracy graph for Hortensia CNN	33
4.2	Accuracy graph for X-Ray CNN	34
4.3	Loss graph for Hortensia CNN	35
4.4	Accuracy graph for X-Ray CNN	35
4.5	Predictions for Hortensia test set	36
4.6	Testing accuracy for Hortensia CNN	37
4.7	Predictions for X-Ray test set	37
4.8	Testing accuracy for X-Ray CNN	38
4.9	Incorrectly classified Hortensia sample	38
4.10	Incorrectly classified X-Ray sample	39

Contents		
		Page No
Acknowledgement		i
Abstract		ii
List Of Tables		iii
List Of Figures		iv
Chapter 1	INTRODUCTION	1
1.1	General Introduction	1
1.2	Present Day Scenario	2
1.3	Motivation	3
1.4	Objective	5
1.5	Schedule	6
1.6	Organisation of Report	7
Chapter 2	BACKGROUND THEORY	8
2.1	Artificial Neural Networks	8
2.2	Representation of images	9
2.3	Convolutional Neural Networks	10
2.4	Loss function and Optimisers	14
2.5	Over-fitting and Under-fitting	15
2.6	Hyper-parameters of Convolutional Neural Networks	16
2.7	Need for Convolutional Neural Networks	17
2.8	Identifying Pneumonia X-Rays	19
Chapter 3	METHODOLOGY	20
3.1	Detailed Methodology	20
3.2	Datasets Used	31
3.3	Software Used	31
Chapter 4	RESULT ANALYSIS	33
4.1	Graphical Results	33
4.2	Testing the Trained CNN	35
4.3	Deviation from Expected Results	38
Chapter 5	CONCLUSION AND FUTURE SCOPE	40
5.1	Conclusion	40
5.2	Future Scope of Work	41

REFERENCES	42
ANNEXURES (OPTIONAL)	43
PROJECT DETAILS	51

CHAPTER 1

INTRODUCTION

This chapter provides an introduction to the basic idea behind the project. It introduces the reader to the concept of big data in healthcare and how the field has rapidly progressed in the recent years. It also gives an insight into the motivation behind the project and the objectives which are intended to be met by it. The schedule for execution of the project can be found at the end of the chapter.

1.1 General Introduction

In today's day and age, there is a vast wealth of digital data that is continuously being generated and stored in various fields. These fields include healthcare, security, e-commerce, social media et cetera. This big data can hold the key to unlocking several patterns and trends if correctly processed. This data exists in several forms. It can be structured alphanumeric data, audio or video data or even images. While the processing of alphanumeric data has progressed by leaps and bounds, processing image type big data is considerably more difficult. This has to do with the fact that computers do not simply see images like humans do. They see images as a matrix of numbers and first and foremost, as seen in Fig. 1.1, before any processing can take place it becomes imperative for the computer to understand what is actually present in the image. This involves the use of image processing tools. If we are to apply artificial intelligence algorithms to image type big data, the sheer dimensionality of the input in image form will lead to very large processing times. Thus, processing image type big data to draw reasonable conclusions is significantly harder than processing alphanumeric data.

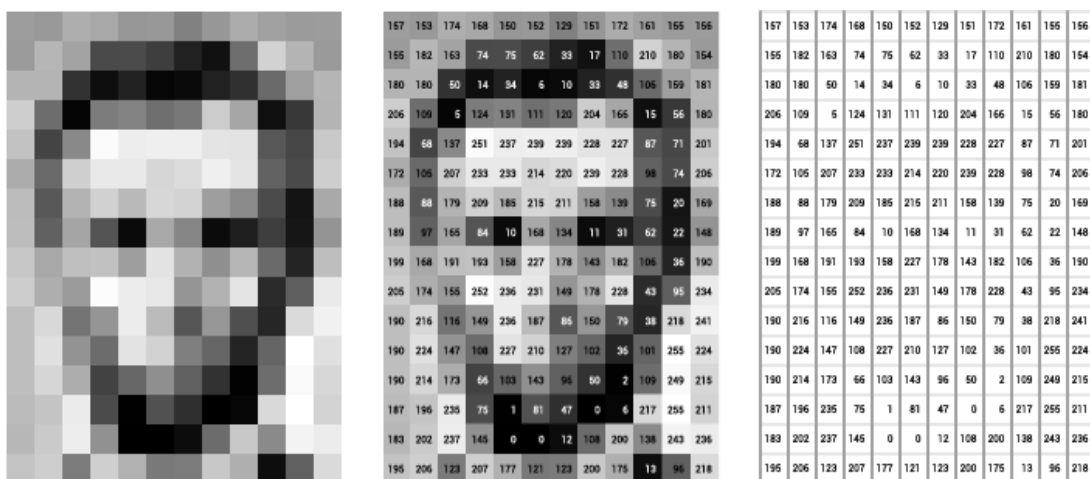


Fig 1.1 Image represented as a matrix

However, a large proportion of the big data being generated is in the form of images. Particularly in the field of medicine, which this project deals with, millions of scans are generated every day. If put to good use, processing these images could lead to new discoveries and trends followed by certain diseases. It could also aid physicians in their diagnosis of diseases. A doctor has a limited experience of having diagnosed a few thousand scans in her lifetime but an artificial intelligence algorithm can be trained to identify a disease on the basis of training it with millions of images, leading to a much higher accuracy than a human physician. This project will use a special type of neural network called a convolutional neural network which is special type of network that can be trained to identify patterns in images. Using a convolutional neural network can do away with several problems that are associated with processing image type big data as it is a network designed especially for image processing.

1.2 Present Day Scenario:

While neural networks have been around for a long times, convolutional neural networks that are a special type of neural networks specifically used to process image type big data are a considerably newer addition. The use of convolutional neural networks has brought about interesting applications like face detection, sentiment analysis and medical image diagnosis today. Although Samsung Research Institute is only just beginning to understand the power of image type big data, several organizations and universities have been working on this specialised technologies for years.

Stanford University has trained powerful convolutional networks to identify if an image of a wart is cancerous or not, as seen in Fig. 1.2.



Fig 1.2 Skin cancer image classification

Its systems can now diagnose skin cancer more accurately than board recognized surgeons. Facebook has made use of the large data it has in terms of pictures of users to train deep convolutional networks to identify people in pictures. Now, its neural networks are so powerful that it can accurately identify one of its users among pictures of over 800 million people. Facebook also identifies objects in its users' pictures to cleverly target advertisements towards them. For example, in Fig. 1.3, mountains are visible, along with a backpack. This will lead to Facebook directing advertisements of backpacks and other trekking gear towards the user.



Fig 1.3 Object recognition by Facebook

Google has trained its convolutional neural networks to perform Super-enhancement, i.e. increasing the quality of poor quality pictures. The quality of a picture is poor because it lacks information. Thus, when you zoom it after a particular stage, large blocks of uniform pixels appear. Google has trained its neural network to identify the closest match out of its dataset to each part of the image and fill in the extra details to enhance quality, as shown in Figure 1.4.

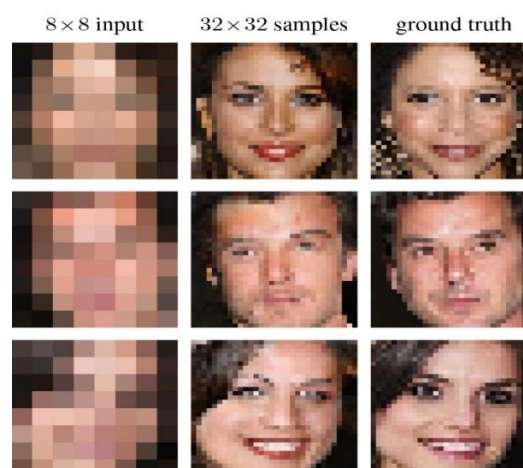


Fig 1.4 Google Super Enhancement

1.3 Motivation

Pneumonia in India accounts for 20% of the deaths due to the disease worldwide. When it comes to children, India has the highest number of pneumonia deaths in the world. UNICEF statistics about pneumonia based deaths of children in India are presented in Table 1.1. While Pneumonia is easy to cure, the deaths are often due to an inability to diagnose the disease from the X-Ray either due to inexperience or negligence. The problem with misdiagnosis of pneumonia is that if the doctor is to diagnose a patient with a less serious condition when it is in fact pneumonia, the infection can progressively get worse and result in serious consequences. On the other hand, if the doctor gives a false positive, when the patient has a different, more serious condition, the results can be disastrous as well. That is why the accurate diagnosis of pneumonia is imperative, especially when one of the most vulnerable sections of population for this disease is young children.

Table 1.1 Pneumonia Death Statistics in Children

Year	% of Neonatal Deaths due to Pneumonia	% of Post- Neonatal Deaths due to Pneumonia	% of Under-5 Deaths due to Pneumonia
2008	6%	30%	17%
2009	6%	30%	17%
2010	6%	30%	17%
2011	6%	30%	17%
2012	5%	31%	17%
2013	5%	31%	16%
2014	5%	30%	16%
2015	5%	28%	15%
2016	5%	28%	15%

This is where a convolutional neural network can easily step in. A neural network can be trained to a sufficiently high level of accuracy wherein it can predict the presence of pneumonia from an X-Ray fed into it without any human intervention. This diagnosis by the network can then be used to aid the doctor to reach a conclusive diagnosis. Samsung is seeking to incorporate artificial intelligence in its digital radiography machines such that a diagnosis can be arrived at, immediately after taking a scan. This diagnosis made by the convolutional neural network (CNN) can be used to aid the doctor or as a standalone as its accuracy increases. Such a system will greatly reduce the chance of misdiagnosis due to negligence or inexperience, as it will be trained on the basis of thousands of images. Such a system will also intimate the patient to seek a second opinion if the diagnosis of the network and the doctor doesn't match.

Incorporating a convolutional neural network in a diagnostic device is a novel idea that hasn't been implemented in any diagnostic devices currently present in India. This is majorly because of the lack of labelled medical datasets that are available openly to train networks on. India has a large wealth of medical data in the form of diagnostic scans taken every day. However, this data mostly serves no purpose except the individuals own, which is a waste of the potential

such sheer volume of data holds. This data can be labelled and used with patients' consent to create large, structured data sets. Samsung is in talks with hospitals in India for generation of elaborate datasets to bring this concept to life. This project uses a publicly available Pneumonia X-Ray dataset.

India has a very low doctor to patient ratio which means a large section of the population has difficulty accessing quality healthcare. Also, due to this lopsided ratio, the time a doctor gets with each patient is significantly limited, which can lead to misdiagnosis. This project hopes to nudge the doctor in the direction of a correct diagnosis even with a short visitation time.

1.4 Objective

This project aims to detect the presence of pneumonia in a lung X-Ray by training a convolutional neural network using a publicly available labelled dataset, for incorporation into digital radiography devices. The network will be able to identify whether the scan is of a healthy person or a person afflicted with Pneumonia. It is proposed to be completed in three phases.

- Phase 1: To understand the limitations of image type big data processing and how a convolutional neural network overcomes these challenges. To learn about the structure, functioning and implementation of a convolutional network, and how to achieve convergence towards a desired result.
- Phase 2: To train a convolutional neural network with simple image data sets like MNIST (handwritten numbers dataset) to understand the architecture and what parameters to tweak to nudge the network towards a higher accuracy. The network should be able to identify the handwritten number provided to it as input (between 0-9) accurately. To then train another CNN with a Hortensia flowers dataset such that the network is able to identify the presence/absence of the flowers in a given input image.
- Phase 3: To train the network with the pneumonia X-Ray dataset and change the number and nature of network layers, and tune the hyper-parameters till a satisfying accuracy is reached and to test the network with new images. The network should be able to tell which scans are afflicted with Pneumonia and which scans are healthy up to a high degree of accuracy.

1.5 Schedule

The tentative schedule for the project is given in Table 1.2.

Table 1.2 Project Schedule

<i>January 2018</i>	<ul style="list-style-type: none">○ Understanding what Big Data is and how it is generated○ Knowledge of the types of big data○ Studying how Big Data analysis is helpful to corporations
<i>February 2018</i>	<ul style="list-style-type: none">○ Research competitor companies usage of Big Data Analytics to make intelligent business decisions○ Study about limitations in image type big data processing
<i>March 2018</i>	<ul style="list-style-type: none">○ Study and understand Convolutional Neural Networks○ Learn about its layers and hyper-parameters○ Learn to use Keras (Neural Network Python Library) and Tensorflow (Google's machine learning architecture)
<i>April 2018</i>	<ul style="list-style-type: none">○ Implement first Convolutional Neural Network○ Train it using MNIST (handwritten characters) dataset○ Tune the hyper-parameters and adjust architecture○ Train another network with more layers for detection of Hortensia flowers
<i>May 2018</i>	<ul style="list-style-type: none">○ Download Pneumonia X-Ray Dataset○ Train a Convolutional Neural Network to detect Pneumonia○ Add layers and adjust nature of layers○ Tune Hyper-parameters till sufficient accuracy is reached○ Check for several test cases
<i>June 2018</i>	<ul style="list-style-type: none">○ Documentation

1.6 Organisation of Report

This report is divided into 5 chapters. The first chapter gives a basic idea into the project, the objective behind it, the tentative schedule and company details. The second chapter, Background Theory, provides a detailed look into the background theory that is necessary to understand the project. The third chapter, Methodology, deals with the approach to the problem and the steps which are planned to execute the project. The fourth chapter deals with the results obtained by following the set plan so far. The fifth and final chapter deals with the conclusions drawn from the obtained results and the steps that can be taken for further improvement.

CHAPTER 2

BACKGROUND THEORY

This chapter deals with the theory required to understand the working of the project. It provides a brief insight to the concept of computer vision and the challenges that are a part of computers identifying the contents of an image. It also provides a detailed understanding of convolutional neural networks. It delves into certain commonly faced issues while training a convolutional neural network. Towards the end of the chapter, there is a discussion on the difference between a healthy lung X-Ray and a lung X-Ray of a patient with Pneumonia.

2.1 Artificial Neural Networks

Artificial neural networks are computational simulations of the networks found in the human brain. Each computational unit in a neural network is called a *neuron*. These neurons are modelled on the behaviour of human brain cells, of the same name. Each neuron can take multiple inputs and provide a single output, as shown in Figure 2.1. The output of a neuron is dependent on the type of *activation function* used, which can be a linear, threshold or sigmoid function. The weighted sum of inputs to the network is given to the function as input. Multiple such neurons come together to form neural networks.

The purpose of a neural network is to allow a system to learn. The network is provided with a large amount of data, called the training dataset. Then it learns the nature of the output desired with each input. This is called *supervised learning*. After learning with the provided dataset, it is tested with a test set to check for the accuracy of what it has learnt. If it has a good accuracy, it can then be used to predict outputs for data it hasn't been trained with. Thus, neural networks allow a system to decide the output on the basis of what it has learnt previously. This is called *soft computing*, as every decision isn't hard coded by the programmer, as seen in hard coding.

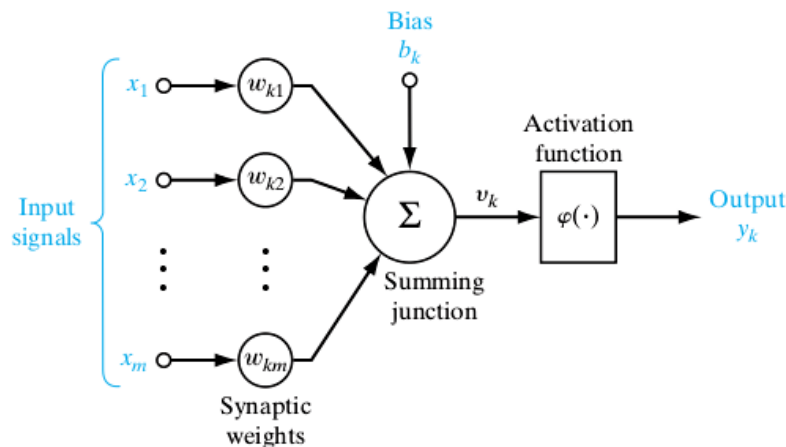


Fig 2.1 Structure of a computational neuron

A neural network can have one or more layers. Usually, for better quality of learning, a neural network has several layers one after the other, such that the output of one layer is fed as input to the consequent layer. Such neural networks are called *multilayer* neural networks.

Feed-forward networks allow the movement of information only in one direction- from input to output. No feedback is permitted. The output of a layer cannot affect that same layer. Such networks are commonly used for *Pattern Recognition*. A neural network in which every neuron of the previous layer is connected to every neuron of the consequent layer, as seen in Figure 2.2, is called a *fully connected neural network*.

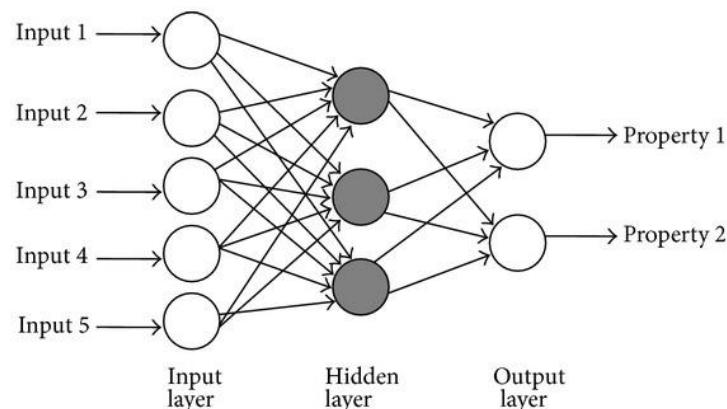


Fig 2.2 A fully connected feed-forward network

2.2 Representation of Images:

Images are represented as a matrix of numbers. For greyscale images, each pixel has a grey level ranging from 0-255. 0 represents the colour black, while 255 is white and all other numbers represent the levels of grey in between. For a colour image, each pixel will consist of three colour channels, red, green and blue. Each of these channels will have a value between 0-255, depending on the intensity of the colour.

Fig. 2.3 shows the split of a colour image into its three constituent colour channels. As each channel has pixels with values ranging from 0-255, each colour channel can easily be shown as a grey scale image with the intensity of white colour corresponding to the intensity of that particular colour channel in the original image.

The representation of images as a matrix of numbers is why computers cannot understand the contents of an image like humans can. This is why they have to be specially trained to see the image as more than a collection of pixel values.

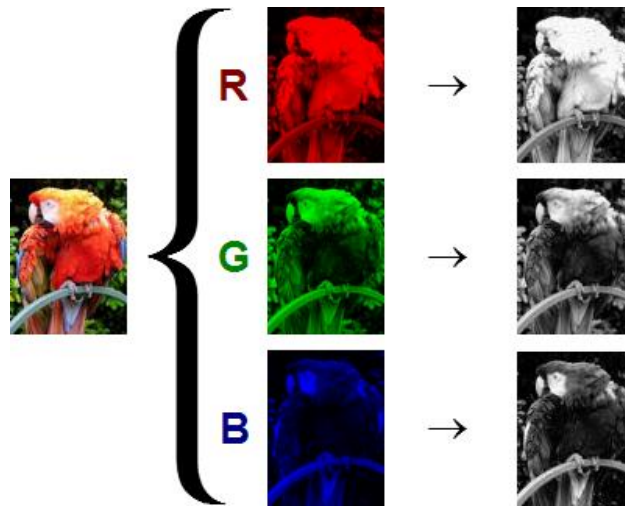


Fig 2.3 RGB splitting of an image

2.3 Convolutional Neural Networks

A convolutional neural network is a specific type of artificial neural network that is specially designed for taking image based inputs. To accommodate this type of input, there are several differences in a convolutional neural network that set it apart from a fully connected, feed forward neural network.

One of the most significant differences is that it is not fully connected. Each node takes inputs from only certain portions of the image, not the entire image itself. Another significant difference is that the weight matrix to each node in a particular node is same. A CNN functions in the form of several layers which will be explained in the subsequent paragraphs. Each of these layers has a specific function and the number, order and placement of these layers can be changed from application to application depending on the nature of the outputs desired. An example of a simple CNN architecture is shown in Fig. 2.4.

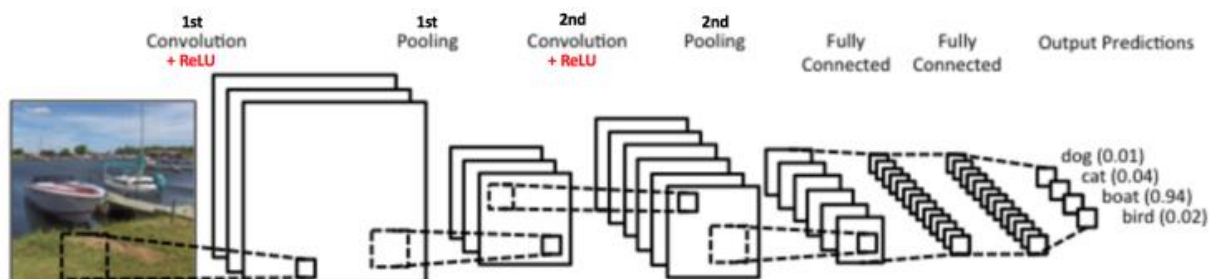


Fig 2.4 Convolutional Neural Network

2.3.1 Convolutional Layer

The most significant layer of a CNN is the convolutional layer, from which the network derives its name. The convolutional layer consists of the following parts:

- **Kernel (The Image Filter):** This is essentially a weight matrix of numbers which glide over the image. The kernel is placed over the top left corner of the image and the pixels corresponding to the kernel are multiplied with the weights of the filter, as shown in Fig. 2.5. The sum of these multiplied weights is stored in the central pixel of the output image. This is called the convolution operation. The nature of the filter decides the kind of output we receive. There are filters to find horizontal and vertical edges, to blur the image etc. In the case of a CNN it starts out with filters of small, positive values, and as the training proceeds, the filters result in clear feature maps of edges and corners.

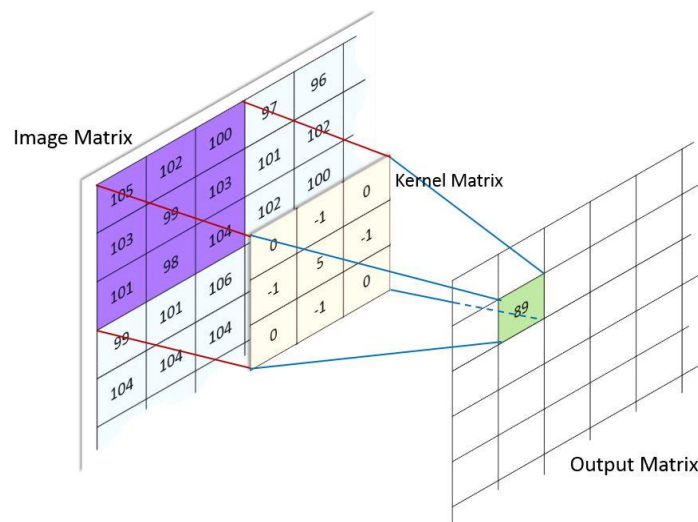


Fig 2.5 A convolution operation on image

- **Sliding and Padding Sizes:** The sliding size determines how far the filter must slide after one computation to perform the subsequent computation. A stride of 1x1 means every slide to the right is of one pixel and every slide to the bottom is also by one pixel value. Sometimes, an image can be padded to allow an output image to be of the same size as the input image. Padding involves adding new pixels all around the boundary of the input image. An example of 0 padding is shown in Fig. 2.6

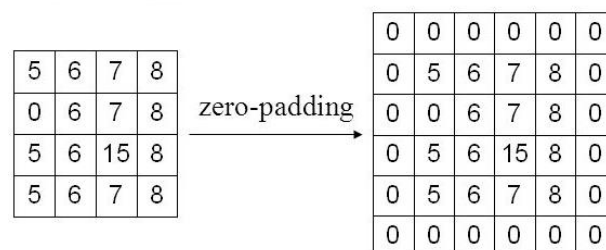


Fig 2.6 Zero Padding an image

- **Feature Maps:** The output image after sliding a kernel all over an input image is called a feature map. It is the convolved output of the image. Each convolutional layer has multiple filters and as each filter results in one feature map, the output of a convolutional layer consists of a stack of feature maps. Fig. 2.7 Illustrates a feature map when a horizontal line kernel is applied on an input image.

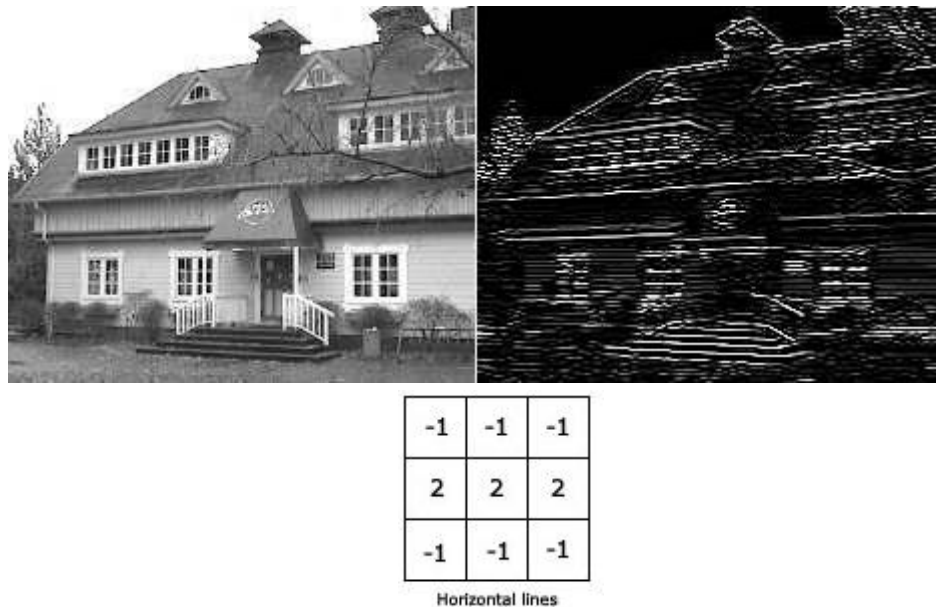


Fig 2.7 (Anti-clockwise from top) a) Input Image b) Kernel c) Feature map

2.3.2 Activation Layer

The activation layer follows a convolutional layer. It takes the feature maps as input to an activation function and applies the activation function to each pixel of each feature map, before passing on the function to the next layer. The choice of the activation function depends on the application, and the position of the activation function in the network. While linear activation functions are preferred inside a network, the final activation function depends entirely on the type of application as explained below.

- **Sigmoid Activation in Binary Classification:** If there are only two possible outcomes, a sigmoid activation function is used. This is essentially a probability function, as seen in Figure 2.8 (a). If the index of one class is 0 and another is 1, an output of 0.5 or less will indicate class 0 and greater than 0.5 will indicate class1.

$$\begin{aligned}
 &\text{if } \frac{e^x}{1 + e^x} < 0.5 \text{ then class 0} \\
 &\quad \text{Otherwise} \quad \quad \quad \text{class 1} \\
 &\text{where } x \text{ is an input to the sigmoid activation function}
 \end{aligned}$$

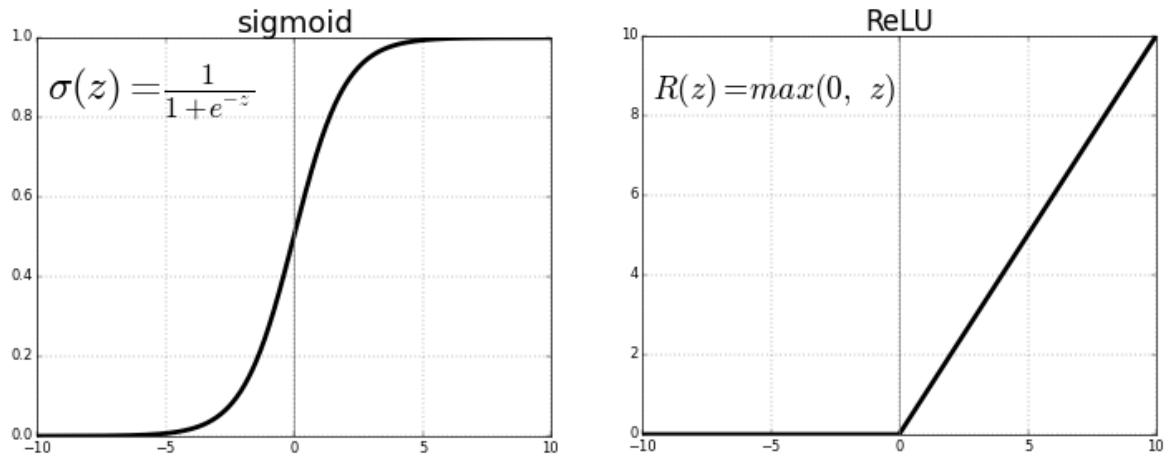


Fig 2.8 a) Sigmoid activation , b) ReLu activation

- Rectified Linear (ReLU) Activation for Linear Regression: If the expected output is a continuous variable, then a ReLU function is used, as seen in Figure 2.8 (b)

$$\begin{aligned} &\text{if } x \geq 0, f(x) = x \\ &\text{else if } x < 0, f(x) = 0 \end{aligned}$$

where x is an input to the ReLU activation function

2.3.3 Pooling Layer

The dimensionality of image data is very large, and as convolutional neural networks have several layers, with image type inputs, the overall computation time would become extremely large, had it not been for pooling layers. Pooling layers essentially work to reduce the dimensionality of the data, by sampling the image data after every few convolutional and activation layers. The major types of pooling layers are explained below.

- Max Pooling: The maximum pixel value out of a group of pixels is sampled and carried on to the downsized data, as seen in Fig. 2.9.
- Average Pooling: The average pixel value out of a group of pixels is sampled and carried on to the downsized data

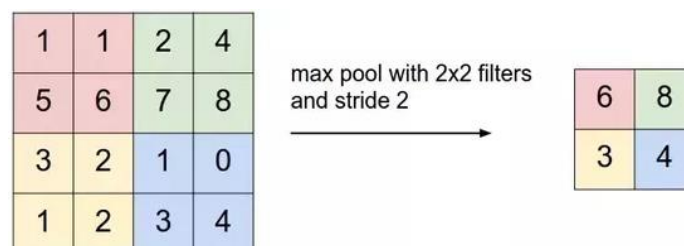


Fig 2.9 Max pooling

2.3.4 Fully Connected Layer

The final layer in a convolutional neural network is a fully connected layer. Each node in this layer takes information from all nodes in the previous layer to make a final prediction. Prior to feeding input to this layer, the output from the previous but one layer, which is in an image format is flattened to a one dimensional array. The activation function applied to the output of a fully connected layer will depend on the application. For the purpose of this project, as binary classification is being performed, a sigmoid function will be chosen for final activation.

2.4 The Loss Function and Optimization

In supervised learning, the CNN knows what output is expected from it for each training image. A loss function is a function that takes the difference between the actual output and the desired output of the CNN as its input. The aim of training a convolutional neural network is to minimise the loss function. The loss function chosen in this project is called binary cross-entropy loss function, as seen in Figure 2.10.

$$\text{loss} = -(y \ln(p) + (1 - y) \log(1 - p))$$

where $y = 0$ or 1 depending on whether the observation matches the desired output, p is the predicted probability of the observation

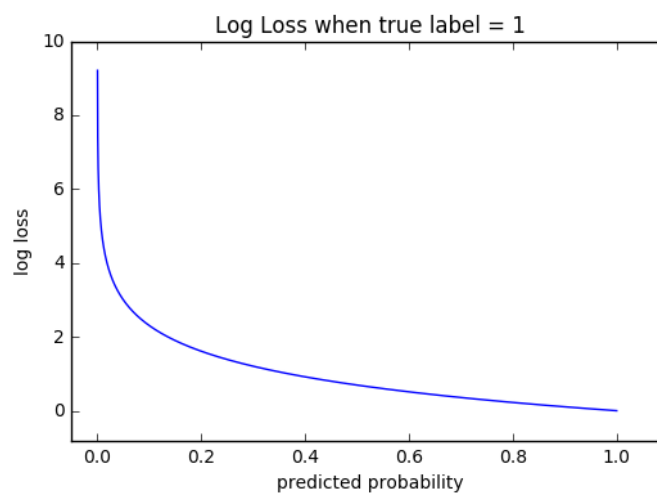


Fig 2.10 Binary cross-entropy loss function

Now, to minimise the loss function, we will differentiate it to find the slope of the loss function and slowly move towards the minimum loss by modifying the weights. This entails taking the gradient of the loss function with each training instance and changing the weights such that we can move downwards along the loss function. This is called gradient descent. When the loss function is minimum, the actual prediction will be closest to the desired outputs. As illustrated

in Fig. 2.11, the weights are modified over the training instances to slowly converge at minimum loss. A generalised understanding of gradient descent is as follows.

$$\begin{aligned} \text{loss} &= f(\text{weight matrix}) \\ \text{delta} &= \text{derivative}(\text{loss}) \\ \text{coefficient} &= \text{coefficient} - (\text{alpha} * \text{delta}), \text{ where alpha is learning rate} \end{aligned}$$

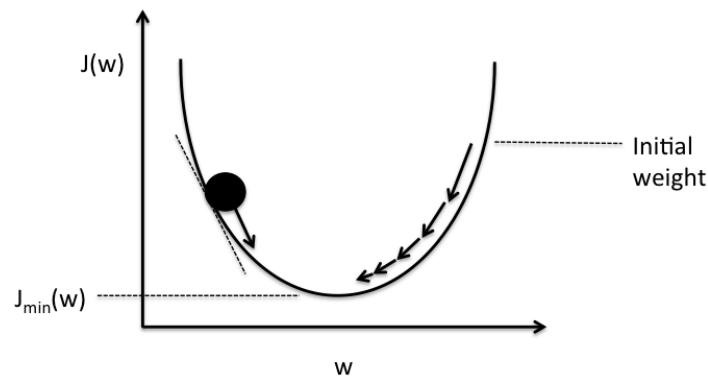


Fig 2.11 Gradient Descent Optimization

2.5 Over-fitting and Under-fitting

While training a neural network, the training data is divided into training and validation data. The training data is used to actually train the network while the validation data is used to intermittently test the network while training it. This is done to prevent phenomena called over-fitting and under-fitting. When the network fits its predictions too closely to the training data, its predictions for data it hasn't encountered before may not be as good as its predictions for the training data.

Validation accuracy is the accuracy of the network's prediction for the validation dataset, and training accuracy is the accuracy of the network's prediction for the training dataset. In the above case, training accuracy will be much greater than the validation accuracy. This is called over-fitting. Over-fitting is a big concern as the accuracy presented while training the model can prove deceptive. One of the ways to combat over-fitting is to include Dropout in the network. Dropout means, in each training cycle, a few nodes selected randomly will be temporarily destroyed, as seen in Figure 2.12. This prevents the network from ever attaining very high training accuracies. Under-fitting is the phenomena opposite to over-fitting when the network fits far too loosely around the training data. In this case, the validation accuracy will be higher than the training accuracy. Watching the growth of the validation and training accuracy simultaneously during training can give a good idea as to whether the model is over or under-fitting and hyper-parameters can be tuned accordingly.

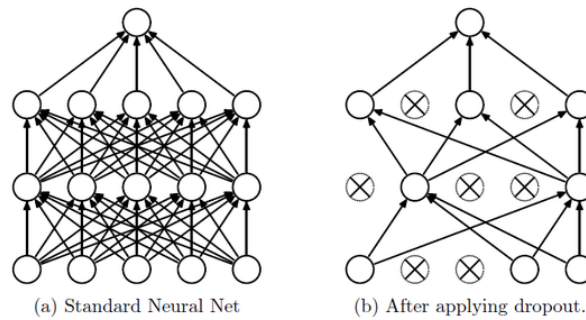


Fig 2.12 Dropout in a neural network

2.6 Hyper-Parameters of a Convolutional Neural Network

There are certain parameters of a convolutional neural network that have to be tuned in accordance to the way the training is proceeding. When the network is not converging towards the desired output, these are the parameters that have to be increased or decreased accordingly.

- **Learning Rate:** As seen in the section 2.5, the learning rate determines by what scale the weights should change so as to move towards minimum loss. A lower learning rate will mean a slower but steadier decline towards minimum loss whereas a higher loss function will mean a speedier yet more haphazard movement towards minimum loss. This can be adjusted according to the desired results.
- **Dropout Rate:** As seen in the previous section, dropout is used to control the problem of over fitting. If a model is over-fitting the dropout rate can be increased and for under-fitting, it can be decreased.
- **Patience:** If a certain metric is not changing after several training instances, the learning rate can be changed automatically in the middle of training. Patience indicates the number of epochs (1 epoch is 1 full training cycle of the training data set) the program must wait before modifying the learning rate.
- **Number and Nature of Layers:** Increasing the number of layers, modifying the number of filters in each convolutional layers, using different types of activation and pooling functions and changing the order of the layers can significantly change the output. If a model is under-fitting, more layers can be added to increase the number of parameters to train. The activation function must be chosen on the basis of the type of application.

2.7 Need for Convolutional Neural Networks

The following are primary issues computers face in understanding images. For a human, an object in an image can be easy to detect in spite of the following, however, for a device that sees images as a matrix of numbers even one of these issues can change the way an image is seen quite significantly.

- **Distortion:** When an object is in a form different from its regular form. In Figure 2.13(a), the cat is in a curled up position that completely changes the angles of its legs and tail from a normal standing cat. This can also be seen if the cat stretches its body or hunches its back.
- **Occlusion:** When an object is partially covered by another object leading to the covering up of important features that lead to the identification of the first object, the first object is said to be occluded by the second. In the Figure 2.13(b), a box is covering the legs and tail of the cat, leading to removal of important identification features.



(a)



(b)

Fig 2.13 (a) Distortion (b) Occlusion

- **Intra-class Variation:** When an object belonging to a particular class can have varying features, it is called intra-class variation. In Figure 2.14(a), there are cats of different shapes and sizes, but all are clearly identifiable as cats.
- **Viewpoint Variation:** The shape and size of the object can vary dramatically on the basis of where you are looking at it from. In Figure 2.14(b), the cat is facing backward, which leads to its facial features being hidden and its tail being seen in prominence.
- **Background Clutter:** When the object to be identified blends in too well with the surroundings, there is said to be background clutter. In Figure 2.14(d), the cat mixes with the floor.

- **Illumination:** Too much or too little lighting can significantly change the appearance of an object. In Figure 2.14(c), the effect of sunlight can make the cat difficult to recognize for a computer.

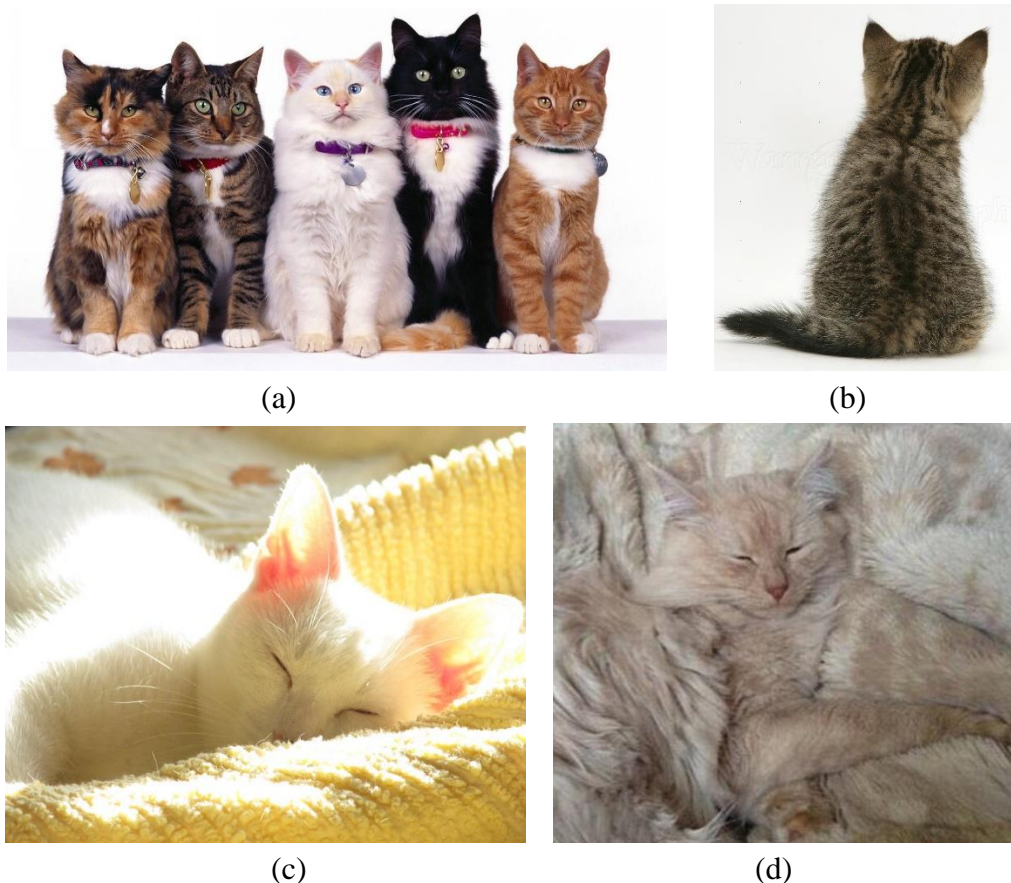
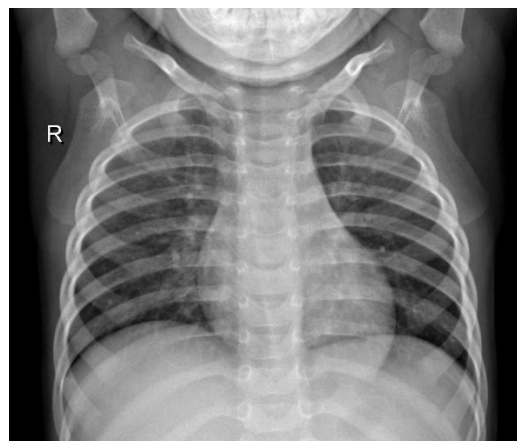


Fig 2.14 (a) Intra-class Variation, (b) Viewpoint Variation, (c) Illumination, (d) Background Clutter

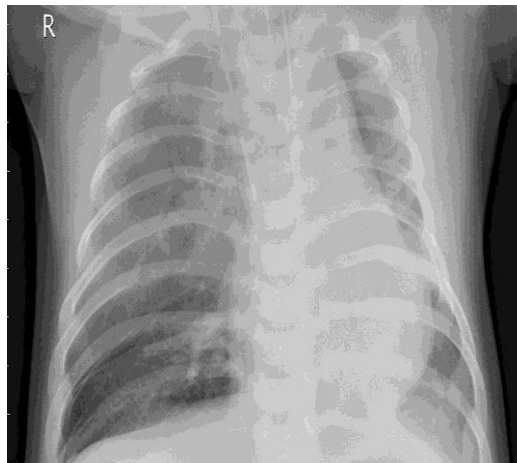
Looking at these issues, it is clear that it is not possible to come up with one single algorithm that can always identify cats. Sometimes some features will be missing, sometimes the cat will be oddly shaped, no two cats look exactly alike et cetera. Thus hard coding is not an option. This is why we will have to find ways to make the system intelligent. Instead of training it to identify a cat by its features, we can show the system pictures of thousands of different cats for it to itself learn what a cat is. This is exactly what deep convolutional neural networks do. We feed several pictures of cats, with the label ‘cat’ to the network and allow it to understand what a cat looks like, as each image passes through all the layers of the CNN.

2.8 Identifying Pneumonia X-Rays

Consolidation is the result of replacement of air in the alveoli by transudate, pus, blood, cells or other substances. Pneumonia is by far the most common cause of consolidation and hence is diagnosed by its presence in X-Rays. The disease usually starts within the alveoli and spreads from one alveolus to another. This consolidation appears as a whitish region in the lungs which should otherwise appear dark. Fig. 2.15 (a) shows an X-Ray of normal lungs. Except for the area occupied by the heart, the lungs are dark in colour and contrast heavily with the ribs. This is an indicator of absence of any consolidation. Fig. 2.15 (b) however shows an X-Ray of Pneumonia infected lungs. The colour of the lungs is much lighter as compared to the healthy lungs and the contrast with the rib cage is less pronounced. This is a result of consolidation.



(a)



(b)

Fig 2.15 (a) Normal lungs X-Ray (b) Pneumonia infected lungs X-Ray

CHAPTER 3

METHODOLOGY

This chapter delves into an in-depth technical explanation of the working of the project. It talks initially about building a simple convolutional neural network using the in-built MNIST data-set. It then discusses developing a more complex network to deal with the X-Ray external data set of higher dimensionality. It discusses the CNN architecture, the parameters set and assumptions made. It also takes a brief look at the tools used to develop the network.

3.1 Detailed Methodology:

This project was implemented in three phases. Initially, a simple convolutional network with very few layers was built to understand the functioning of CNN's, tuning of hyper-parameters and the arrangement of layers, followed by a more complex CNN for Pneumonia detection.

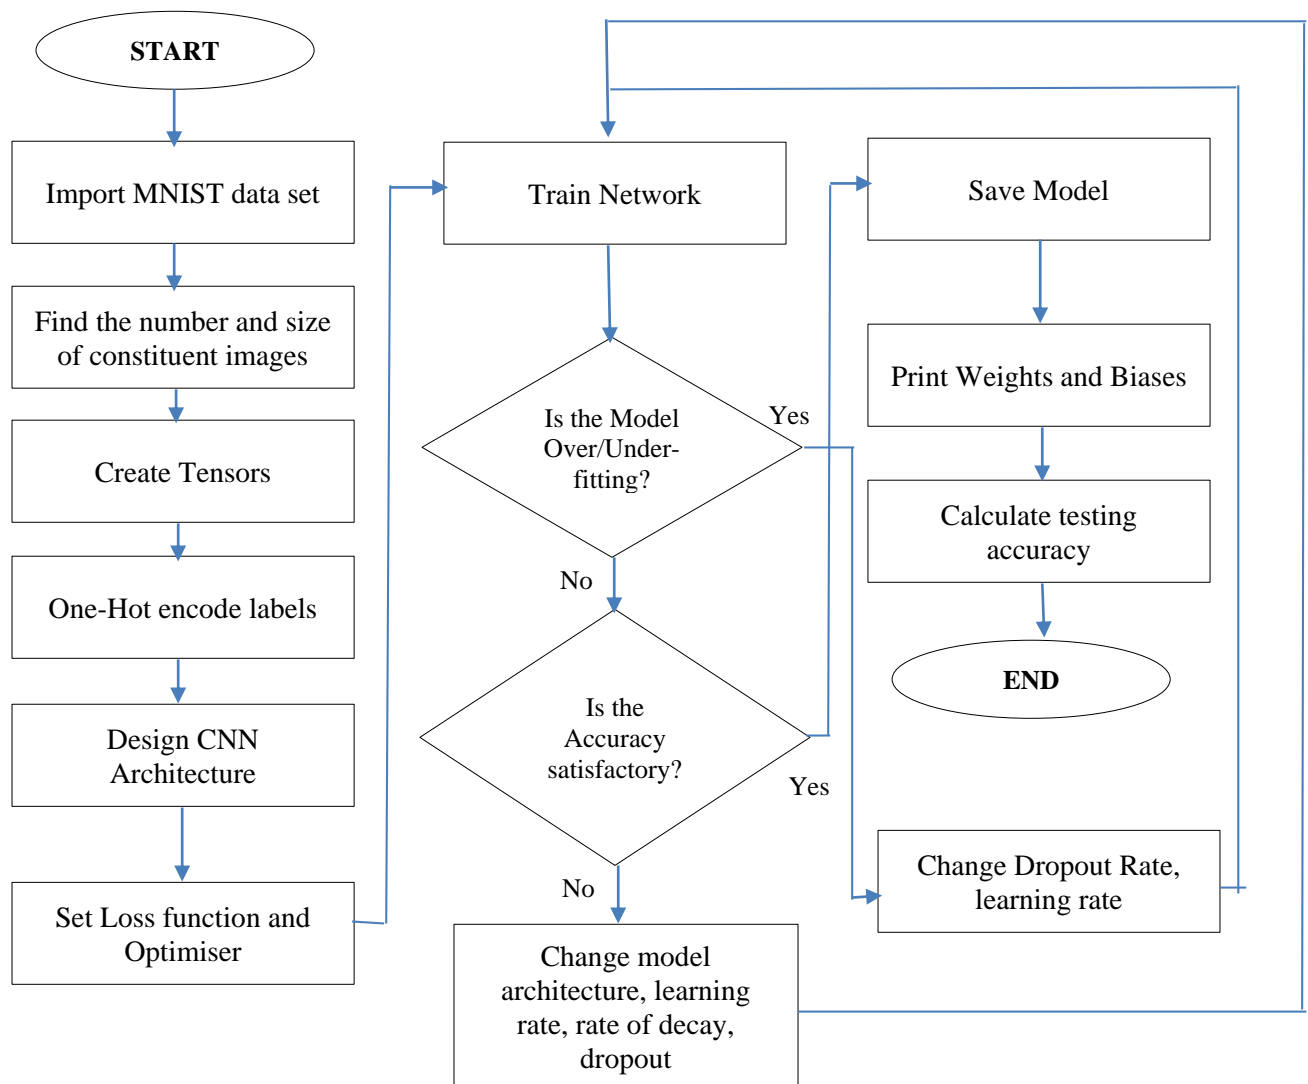


Fig 3.1 Flowchart of MNIST CNN Methodology

3.1.1 MNIST CNN Methodology

To start the project, understanding the functioning of a CNN was necessary. For this, the initial CNN was built with a small number of layers to train it to recognize handwritten characters from the MNIST data-set. The Methodology is displayed in Fig. 3.1.

- Import MNIST data-set: MNIST data-set is an elementary data set used to understand the working of machine learning algorithms. It consist of 70,000 images of handwritten numbers between 0-9 of the image size 28x28x1, as seen in an example in Figure 3.2. This dataset is available in Keras, the machine learning library of python. Thus, it was loaded directly from there.

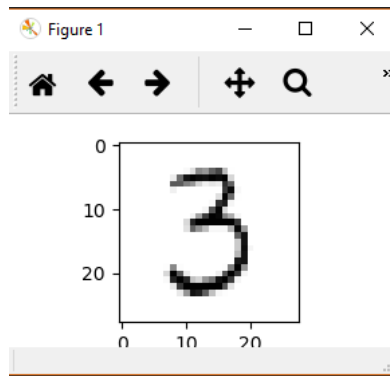


Fig 3.2 MNIST image example

- Find number of Images and shape of Images: Initially the MNIST dataset consists of 60000 training and 10000 testing images, as seen in Figure 3.3 (a). The training dataset is further divided into training and Validation data-sets and the number of images in each of the sets is counted, as seen in Figure 3.3(b). The shape of the images are confirmed as 28x28x1 as seen in Figure 3.2.

```
train_images size (60000, 28, 28)
train_labels size (60000,)
test_images size (10000, 28, 28)
train_labels size (10000,)
```

(a)

```
train images size (55000, 28, 28)
validation images size (5000, 28, 28)

Process finished with exit code 0
```

(b)

Fig 3.3 (a) Test and Training dataset sizes (b) Size after division of training set

- Create Tensors: In this project, Google's tensor flow, which is a neural network infrastructure is being used. This structure requires the inputs to be in the form of tensors, with all pixel values lying between 0-1. Thus, all pixel values are divided by 255 and converted into a float data type (for decimal values).
- One hot encode labels: Each of the images in MNIST is labelled as a number between 0-9. However, processing in Tensor-flow becomes simplified if the labels are one-hot encoded. This means if a label is 4, a ten unit one dimensional array will be created with all indices except the 4th index as 0, and the 4th index as 1, i.e. [0 0 0 0 1 0 0 0 0 0]
- Design CNN Architecture: A tentative CNN architecture is designed which can be tweaked as per conditions later. A CNN architecture entails defining the layers and the parameters within each layer. In this example, all activations layers will make use of the ReLu activation function, except the final layer, which will use a sigmoid activation function. The model summary, i.e. a detailed list of all the layers in the final CNN design is shown in Figure 3.4.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 16)	160
activation_1 (Activation)	(None, 26, 26, 16)	0
conv2d_2 (Conv2D)	(None, 24, 24, 16)	2320
activation_2 (Activation)	(None, 24, 24, 16)	0
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 16)	0
dropout_1 (Dropout)	(None, 12, 12, 16)	0
conv2d_3 (Conv2D)	(None, 10, 10, 8)	1160
activation_3 (Activation)	(None, 10, 10, 8)	0
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 8)	0
flatten_1 (Flatten)	(None, 200)	0
dropout_2 (Dropout)	(None, 200)	0
dense_1 (Dense)	(None, 10)	2010
activation_4 (Activation)	(None, 10)	0
dense_2 (Dense)	(None, 10)	110
activation_5 (Activation)	(None, 10)	0
Total params: 5,760		

Fig 3.4 Model summary of MNIST CNN

- The loss function and optimiser type are chosen. Owing to the simplicity of the dataset, the most basic optimiser stochastic gradient descent is chosen. The loss function is categorical cross entropy, owing to the multiple classes.
- Training the Network: The number of epochs for which the training will continue are chosen. Each epoch means 1 cycle of the entire training data passing through the CNN. As this was just to understand the functioning, 20 epochs were selected. The training is done in batches in each epoch to reduce computation time. 2500 instances were selected to be in each batch.

$$\text{Steps per epoch} = \frac{\text{No. of trainig images}}{\text{Batch Size}} = \frac{55000}{2500} = 22$$

- Under-fitting was observed so the dropout rate was changed from 0.4 to 0.15 in all dropout layers.
- Saving the model: The fully trained model was saved so that it could be directly loaded next time without having to train it again to make predictions.
- Printing weights and biases: This is an optional step, if the weights and biases corresponding to any layers are required to be inspected. It displays the numerical value of the weights pertaining to a certain layer.
- Checking Test Set Accuracy: The trained model was initially tested for individual images, as shown in Fig. 3.5 from the test set and then the accuracy of prediction of the model was tested with the entire test set as a whole.

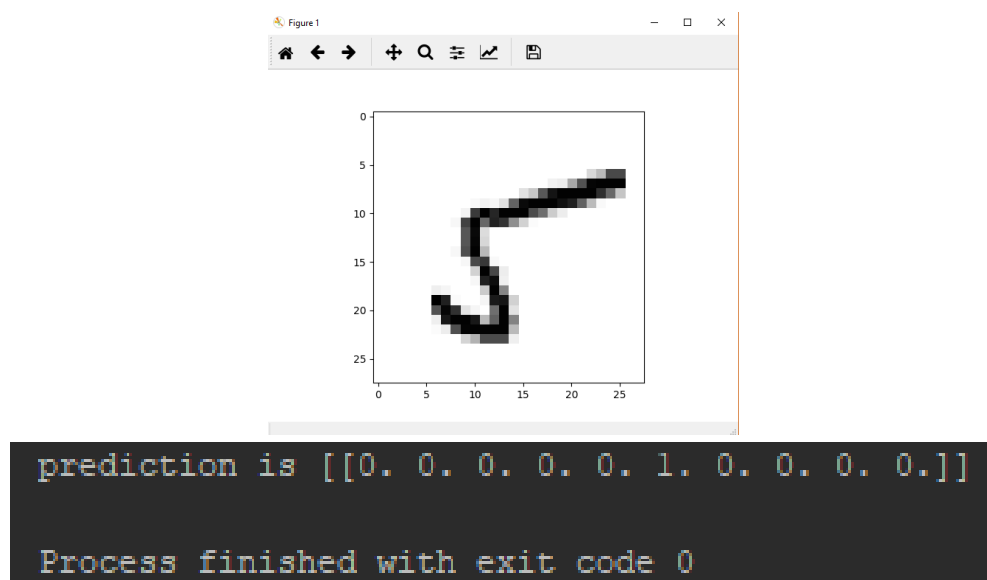


Fig 3.5 Testing the MNIST CNN

3.1.2 X-Ray CNN Methodology

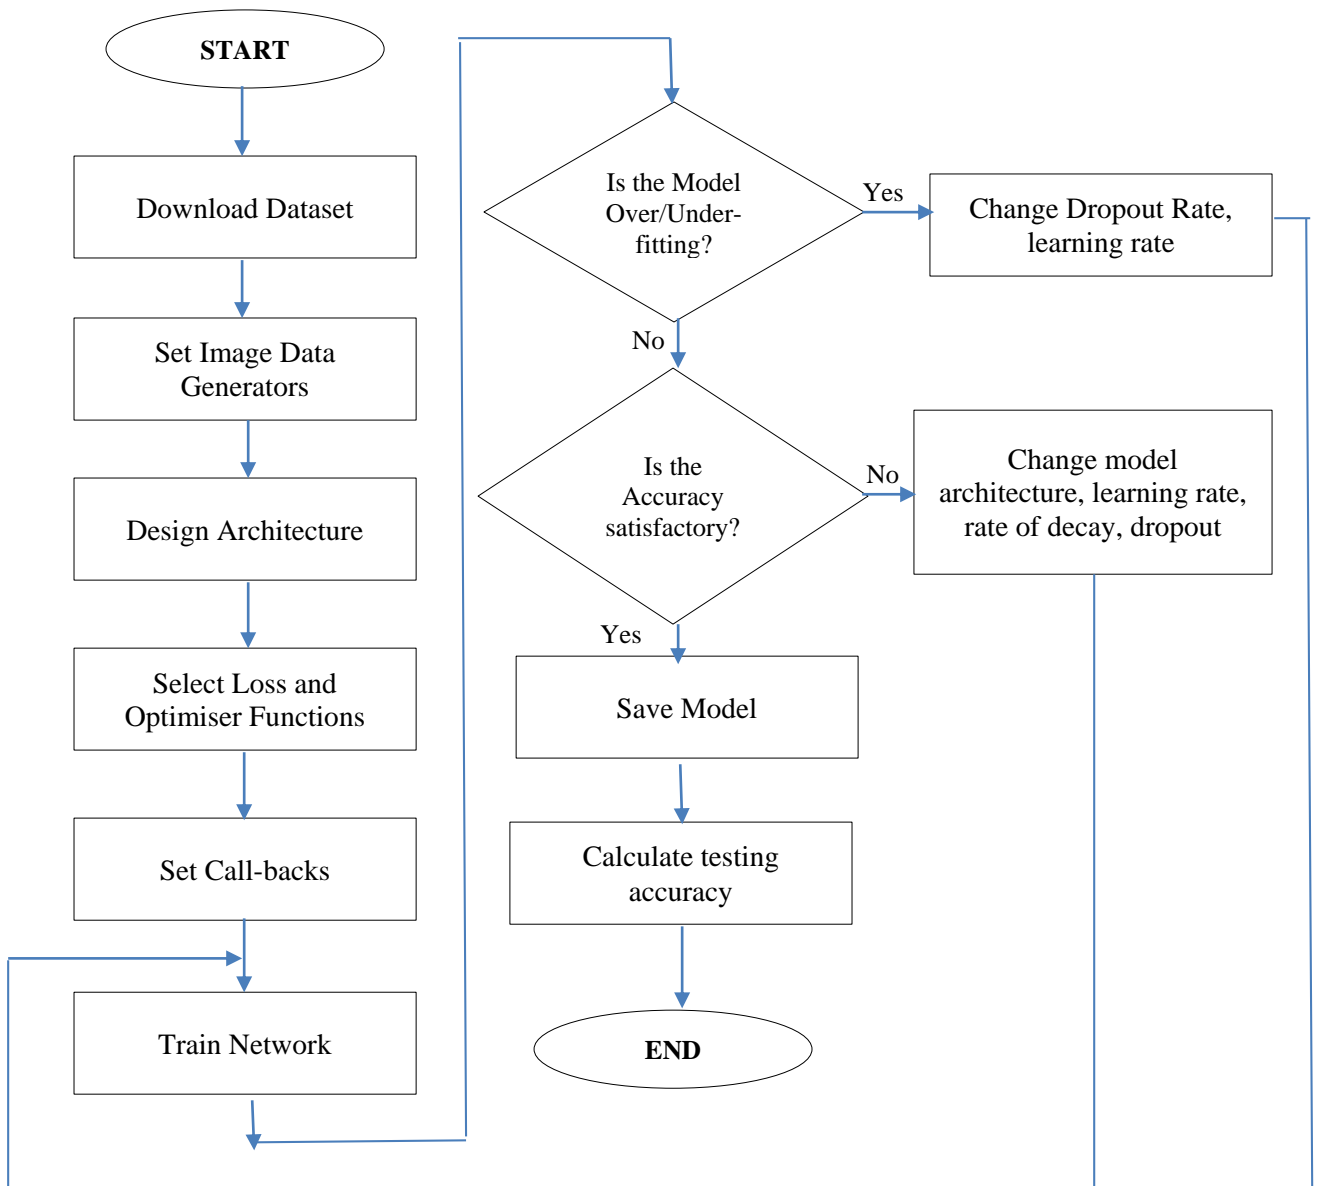


Fig 3.6 Flowchart representation of X-Ray CNN Methodology

Training a CNN for the Hortensia flowers data-set and the pneumonia X-Ray Dataset use the same methodology, with minor differences from the MNIST data-set training. These differences arise from the fact that the image dimensionality is larger for these datasets. Also, we are training these CNN's for binary classification unlike in the MNIST case with multi-class classification. Another important factor is the difference between the classes is more subtle in the case of Hortensia dataset and X-Ray data-set than in the case of MNIST where all classes are very distinguishable from each other. Another difference is that Hortensia data-set and X-Ray data-set are downloaded externally, not from Keras.

The methodology for the training CNN for Hortensia and X-Ray Dataset is explained below:

- The aim of training the CNN with Hortensia data is for the network to correctly predict the presence or absence of Hortensia flowers in an input image. For the X-Ray data, the network should be able to decide whether the scan is of a healthy person or a person with Pneumonia.
- The datasets were downloaded from Kaggle.com. The Hortensia dataset contained 2849 images in the training set, 500 images in the validation set and 20 images in the testing set, as seen in Fig 3.7. The X-Ray Dataset consisted of 5216 training images, 624 validation images, and 16 testing images, as seen in image 3.8.

```
Size of Hortensia Dataset
No. of images in Train data-set= 2849
No. of images in Test data-set= 20
No. of images in Valid data-set= 500

Process finished with exit code 0
```

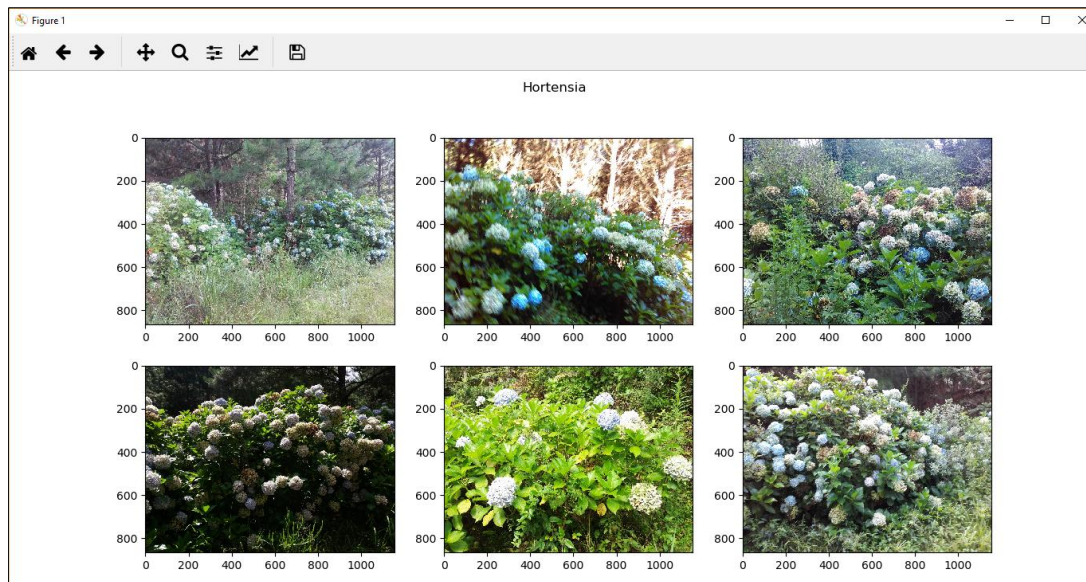
Fig 3.7 Hortensia data-set size

```
X-RAY data-set size
No. of images in Train data-set= 5216
No. of images in Test data-set= 16
No. of images in Valid data-set= 624

Process finished with exit code 0
```

Fig 3.8 Size of X-Ray Dataset

- Now to understand the types of images we are dealing with, and the image sizes, we randomly sample 6 images each from the training sets of Hortensia flowers, i.e. 6 images with Hortensia flowers and 6 images without, as displayed in Fig. 3.9. The size of all the images is the same which is 866x1154x3, as shown in Fig. 3.10. The third dimension indicates that the images are colour images with RGB channels. Similarly before training the X-Ray CNN, 6 Samples each of healthy and Pneumonia affected scans were randomly sampled and printed, as seen in Fig. 3.11. The size of images in this case isn't the same as can be seen in Fig. 3.12. This will later be addressed when all images are resized to the same size.



(a)

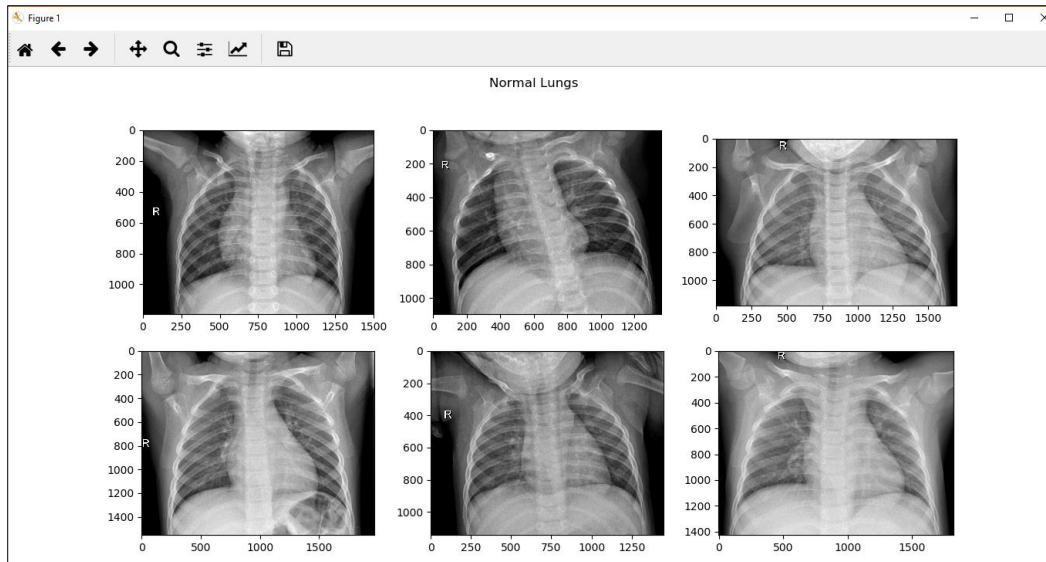


(b)

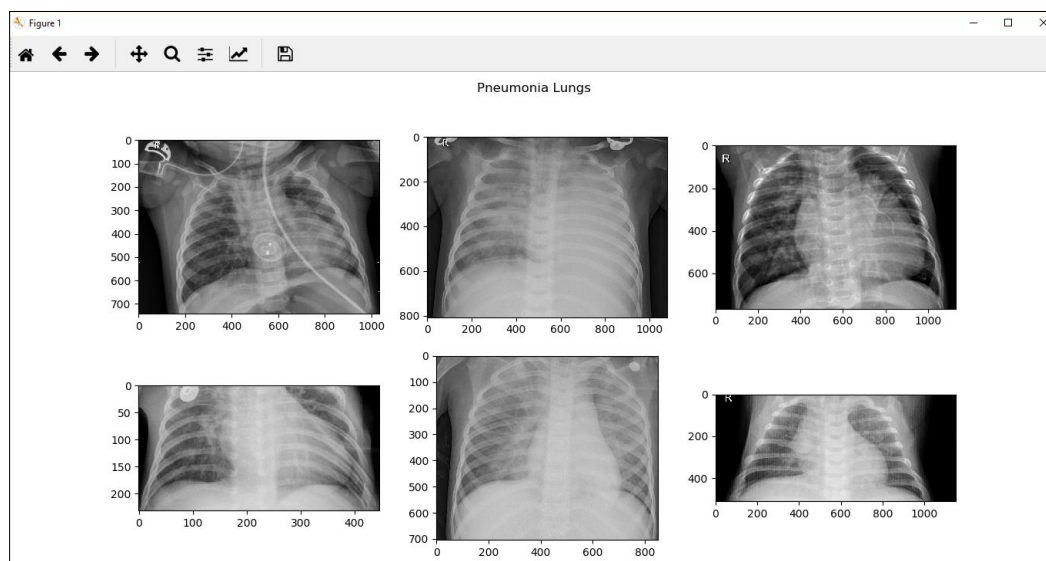
Fig 3.9 (a) Hortensia Flower Images (b) Without Hortensia Images

```
Hortensia Samples
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
No Hortensia Samples
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
(866, 1154, 3)
Process finished with exit code 0
```

Fig 3.10 Sample Hortensia image sizes



(a)



(b)

Fig 3.11 (a) Healthy X-Ray samples (b) Pneumonia X-Ray samples

```

Normal Lung Samples
(1194, 1500)
(1097, 1360)
(1178, 1700)
(1554, 1970)
(1144, 1450)
(1426, 1824)
Pneumonia Lung Samples
(744, 1032)
(808, 1080)
(766, 1132)
(232, 446, 3)
(704, 851, 3)
(512, 1152)

```

Fig 3.12 X-Ray Sample sizes

- **Image Data Generators:** When the training data is of a small size, as in a few thousand samples, it is necessary to introduce some level of variation in the training data before each epoch to avoid over-fitting. The data must also be scaled down to pixel values between 0 and 1 for processing in Tensor Flow. This is done by image data generators. Image data generators allow us to set the height shift range, width shift range, zoom range and rescale parameters to introduce variation in the training images. As the Hortensia image dataset could be easily distinguished into two classes, variation parameters could be set without resulting in under-fitting. Thus all above parameters were used to introduce variation. However, as the X-Ray images had more subtle differences between the classes, too much variation by image data generators could result in under-fitting, thus only rescale parameter was used to divide pixel values by 255.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 99, 99, 16)	208
activation_1 (Activation)	(None, 99, 99, 16)	0
max_pooling2d_1 (MaxPooling2)	(None, 49, 49, 16)	0
dropout_1 (Dropout)	(None, 49, 49, 16)	0
conv2d_2 (Conv2D)	(None, 48, 48, 16)	1040
activation_2 (Activation)	(None, 48, 48, 16)	0
conv2d_3 (Conv2D)	(None, 47, 47, 16)	1040
activation_3 (Activation)	(None, 47, 47, 16)	0
max_pooling2d_2 (MaxPooling2)	(None, 23, 23, 16)	0
dropout_2 (Dropout)	(None, 23, 23, 16)	0
conv2d_4 (Conv2D)	(None, 22, 22, 16)	1040
activation_4 (Activation)	(None, 22, 22, 16)	0
max_pooling2d_3 (MaxPooling2)	(None, 11, 11, 16)	0
conv2d_5 (Conv2D)	(None, 10, 10, 16)	1040
activation_5 (Activation)	(None, 10, 10, 16)	0
max_pooling2d_4 (MaxPooling2)	(None, 5, 5, 16)	0
flatten_1 (Flatten)	(None, 400)	0
dropout_3 (Dropout)	(None, 400)	0
dense_1 (Dense)	(None, 10)	4010
activation_6 (Activation)	(None, 10)	0
dense_2 (Dense)	(None, 1)	11
activation_7 (Activation)	(None, 1)	0
Total params: 8,389		
Trainable params: 8,389		
Non-trainable params: 0		

Fig 3.13 Hortensia CNN Model Summary

- The architecture, in terms of the number and nature of layers was set after trying several architectural styles during training. In case of the Hortensia CNN, fewer number of layers sufficed, but the X-Ray dataset owing to more subtle differences between classes required a large number of layers. The final architecture of the Hortensia CNN arrived at is displayed in Fig 3.13 and the final CNN architecture of the X-Ray dataset is shown in Fig. 3.14. In the final architecture of Hortensia CNN, the activation layers are displayed separately, whereas in the X-Ray CNN, the activation layers are incorporated in each Convolution layer itself. Both methods are functionally equivalent, but this was done to minimize the appearance of the model summary of the X-Ray CNN which would otherwise appear much larger.

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 64)	1792
conv2d_2 (Conv2D)	(None, 146, 146, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 64)	0
dropout_1 (Dropout)	(None, 73, 73, 64)	0
conv2d_3 (Conv2D)	(None, 71, 71, 32)	18464
conv2d_4 (Conv2D)	(None, 69, 69, 32)	9248
max_pooling2d_2 (MaxPooling2D)	(None, 34, 34, 32)	0
dropout_2 (Dropout)	(None, 34, 34, 32)	0
conv2d_5 (Conv2D)	(None, 32, 32, 16)	4624
conv2d_6 (Conv2D)	(None, 30, 30, 16)	2320
max_pooling2d_3 (MaxPooling2D)	(None, 15, 15, 16)	0
dropout_3 (Dropout)	(None, 15, 15, 16)	0
flatten_1 (Flatten)	(None, 3600)	0
dense_1 (Dense)	(None, 128)	460928
dropout_4 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 1)	129
Total params: 534,433		
Trainable params: 534,433		
Non-trainable params: 0		

Fig 3.14 X-Ray CNN model summary

- The loss function and optimiser type are chosen. For the Hortensia CNN, stochastic gradient descent optimiser is chosen. For X-Rays, a slightly more sophisticated form of the gradient descent optimiser called ‘Adam’ is chosen. As both problems deal with binary classification, a binary cross-entropy function is chosen to be the loss function.
- Setting Call-backs: Two call-backs are set. The first call back is called a model checkpoint. This allows the program to keep saving the model with the highest validation accuracy after each epoch, so that if the training were to be stopped mid-way, the best model that had been trained up till that point is still accessible. The second call-back is used to reduce the learning rate if the validation accuracy remains constant for a certain number of epochs. This number of epochs for which the program waits before reducing the learning rate is called patience. As the Hortensia CNN had a larger no. of epochs, its patience was set at 5, whereas for the x-ray CNN, patience was 2. The decrease in learning rate after plateau of validation accuracy for 5 epochs is shown in Fig. 3.15.

```

19/22 [=====>.....] - ETA: 10s - loss: 0.6239 - acc: 0.6574
20/22 [=====>...] - ETA: 6s - loss: 0.6260 - acc: 0.6549
21/22 [=====>...] - ETA: 3s - loss: 0.6263 - acc: 0.6576
22/22 [=====] - 95s 4s/step - loss: 0.6253 - acc: 0.6601 - val_loss: 0.6467 -
val_acc: 0.5640

Epoch 00017: val_acc did not improve from 0.73000
Epoch 00017: ReduceLROnPlateau reducing learning rate to 0.14250000566244125.
Epoch 18/100

```

Fig 3.15 Change of learning rate while training

- Training the Network: The number of epochs for which the training will continue are chosen. Each epoch means 1 cycle of the entire training data passing through the CNN. The Hortensia CNN was trained for 100 epochs as it took longer to converge owing to a simpler architecture. The X-Ray CNN was trained for 12 epochs, owing to a more complex architecture.
- Saving the model: The fully trained models were saved so that they could be directly loaded next time without having to train them again to make predictions. The models with the highest validation accuracy were used thereafter.
- Checking Test Set Accuracy: The trained models were initially tested for individual images, the test set and then the accuracy of prediction of the model was tested with the entire test set as a whole.

3.2 Datasets Used

The following datasets were used for developing this project.

- MNIST Handwritten Characters Dataset – Inbuilt dataset in Keras (Machine Learning Library) of Python. It consist of 70,000 images of handwritten digits, of size 28x28x1.
- The Hortensia Dataset – This dataset consisted of images that either had Hortensia flowers in them or didn't. All images were of the same size. The dataset was downloaded from <https://www.kaggle.com/alxmamaev/flowers-recognition>
- The Pneumonia X-Ray Dataset: Chest X-ray images (anterior-posterior) were selected from retrospective cohorts of pediatric patients of one to five years old from Guangzhou Women and Children's Medical Center, Guangzhou. All chest X-ray imaging was performed as part of patients' routine clinical care. For the analysis of chest x-ray images, all chest radiographs were initially screened for quality control by removing all low quality or unreadable scans. The diagnoses for the images were then graded by two expert physicians before being cleared for training the AI system. In order to account for any grading errors, the evaluation set was also checked by a third expert. The dataset was downloaded from <https://data.mendeley.com/datasets/rscbjbr9sj/2>

3.3 Software Used

This project was executed using Python version 3.6.5. This project made extensive use of the following:

- Keras Library: Keras is a high-level deep learning library, written in Python and capable of running on top of TensorFlow, CNTK, or Theano, which are different computation platforms. It was developed with a focus on enabling fast experimentation. It has a high degree of modularity. In particular, neural layers, cost functions, optimizers, initialization schemes, activation functions, regularization schemes are all standalone modules that you can combine to create new models.
- TensorFlow: Its flexible architecture allows easy deployment of computation across a variety of platforms (CPUs, GPUs, TPUs), and from desktops to clusters of servers to mobile and edge devices. Originally developed by researchers and engineers from the Google Brain team within Google's AI organization, it comes with strong support for machine learning and deep learning and the flexible numerical computation core is used across many other scientific domains.

This chapter went into detail about how a convolutional network can be trained to fit itself to a certain dataset well enough that when new input images are provided to it, it can make an accurate prediction. It started out building a network with a simple MNIST dataset, proceeded through the Hortensia dataset and finally trained a network to do what we had started out to, predict the presence of Pneumonia in X-Ray inputs.

CHAPTER 4

RESULT ANALYSIS

This chapter analyses the performance of the Convolutional Neural Networks using the Hortensia data-set and the Pneumonia X-Ray Dataset. It evaluates the growth of the validation and training accuracy while the training is going on. It later has a look at how well the network is able to predict the output for images it hasn't been exposed to during training, from the testing dataset.

4.1 Graphical Results

While training a CNN, it is pivotal to watch the growth of the training and validation accuracy and the decrease of the training and validation loss with the no. of epochs to train it better, as explained in the following sub-sections. This helps to see if training is proceeding in the right direction, and whether there is any under/over-fitting take place.

4.1.1 Variation of Training and Validation Accuracy with No. of Epochs

The variation of training and validation accuracy with the no. of epochs for the CNN trained with the Hortensia data-set is shown in Figure 4.1. Validation accuracy is the accuracy of the predictions made by the CNN for all images in the validation dataset. The validation accuracy is checked at the end of each epoch and compared with the training accuracy to check for any over-fitting or under-fitting. Training accuracy is continuously measured after each batch of images within each epoch. Training accuracy is the accuracy of predictions by the CNN for all images in the training dataset. As seen in Fig 4.1, the training and validation accuracy grow together, which means there isn't under-fitting or over-fitting going on.

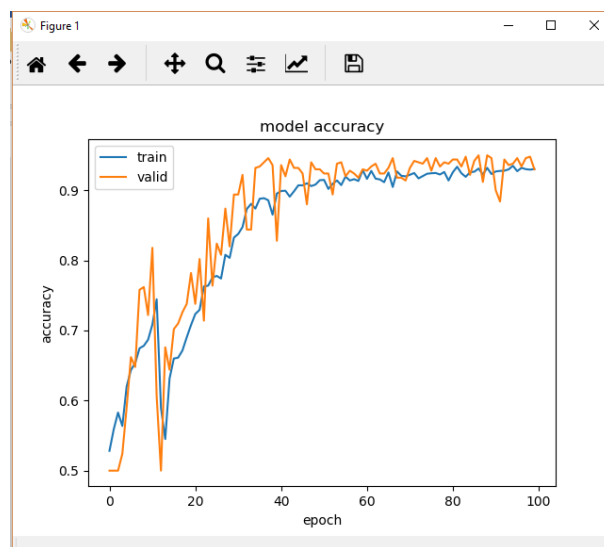


Fig 4.1 Accuracy Graph for Hortensia CNN

The validation accuracy reaches a maximum value of 95% at the end of 87 epochs and then again decreases to 93% at the end of 100 epochs. At the end of 100 epochs, the training accuracy is 93%. Using the call-backs, the model with the highest validation accuracy at 95% is then used for testing. Fig. 4.2 displays the validation and training accuracy vs. no. of epochs for the CNN trained with the Pneumonia X-ray dataset. Owing to the larger number of parameters in this CNN architecture, it is trained for only 13 epochs. The validation accuracy reaches a maximum value of 90% at the end of 11 epochs and then again decreases to 89.5% at the end of 13 epochs. At the end of 13 epochs, the training accuracy is 93%. Using the call-backs, the model with the highest validation accuracy at 90% is then used for testing. In this graph, it is noticeable that the training accuracy is always slightly greater than the validation accuracy. Thus, there is mild over-fitting. However, owing to the fact that the differences between the two classes are very subtle in this binary classification problem, we need a larger number of parameters in the CNN to reach a high degree of accuracy. When the number of parameters increases, a mild level of over-fitting is unavoidable.

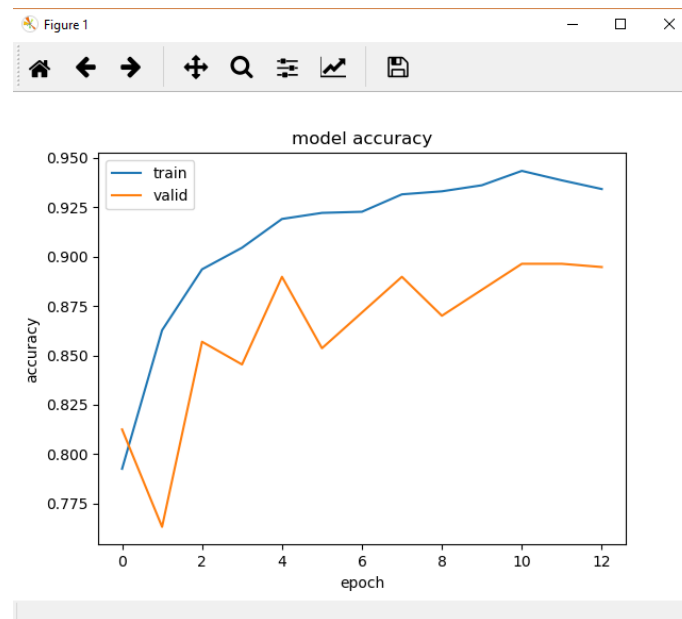


Fig 4.2 Accuracy Graph for X-Ray CNN

4.1.2 Variation of Training and Validation Loss with No. of Epochs

An increase in accuracy is marked by a decrease in the loss function. A loss function is a function that takes the difference between the actual prediction and the desired prediction as its input. Higher the difference, higher the loss function, thus a decrease in the loss indicates training proceeding in the right direction. The loss function used is binary cross-entropy. In the case of Fig. 4.3, the variation in training and validation loss with no. of epochs for the Hortensia dataset trained CNN is shown. This graph confirms the implications derived from the accuracy graph. Both loss decrease at a similar rate and very close to each other, indicating the absence of over/under-fitting. Fig 4.4 shows the loss graph for the X-Ray dataset trained CNN.

As seen in the accuracy graph, there is mild over-fitting which means validation loss decreases less and more slowly than the training loss. However, there is a general trend of decrease of loss which means the network is training in the right direction.

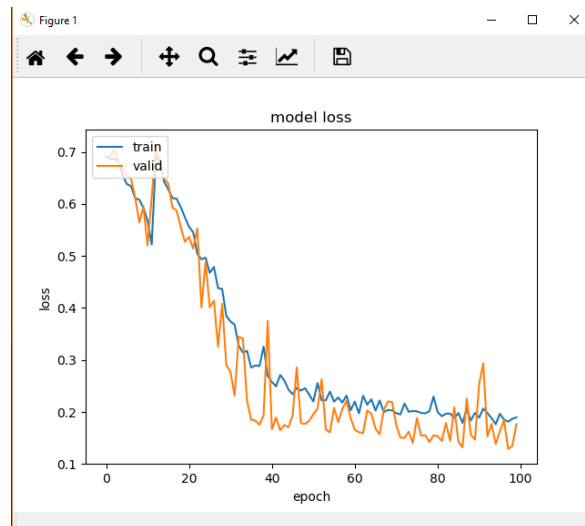


Fig 4.3 Loss Graph for X-Ray CNN

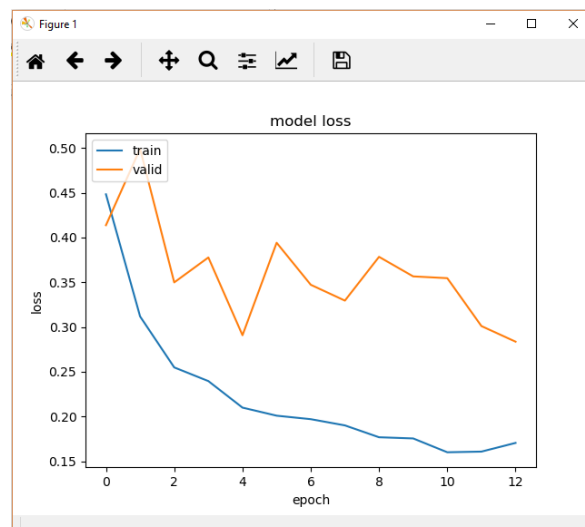
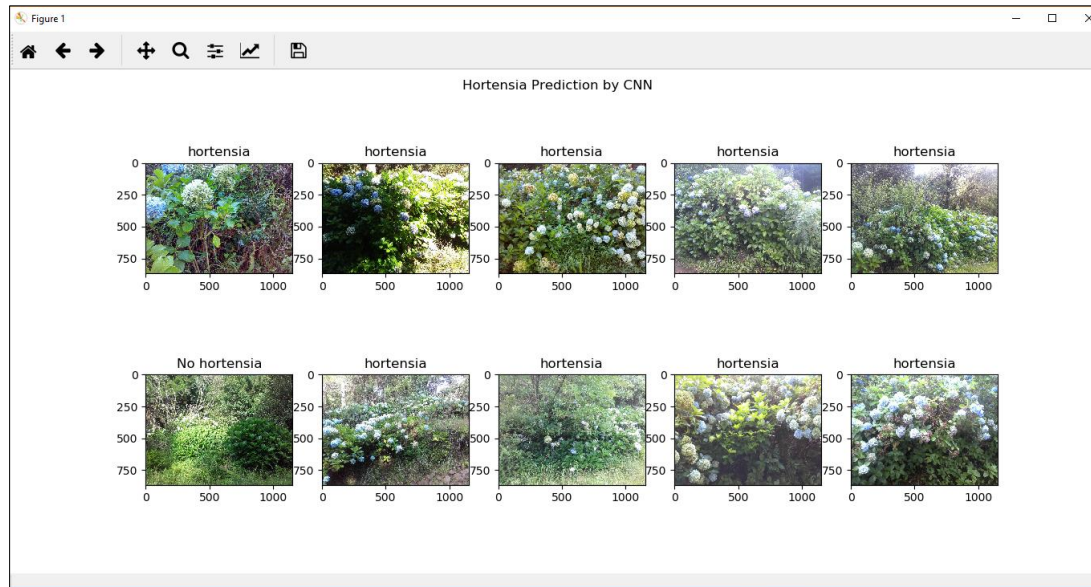


Fig 4.4 Loss Graph for X-Ray CNN

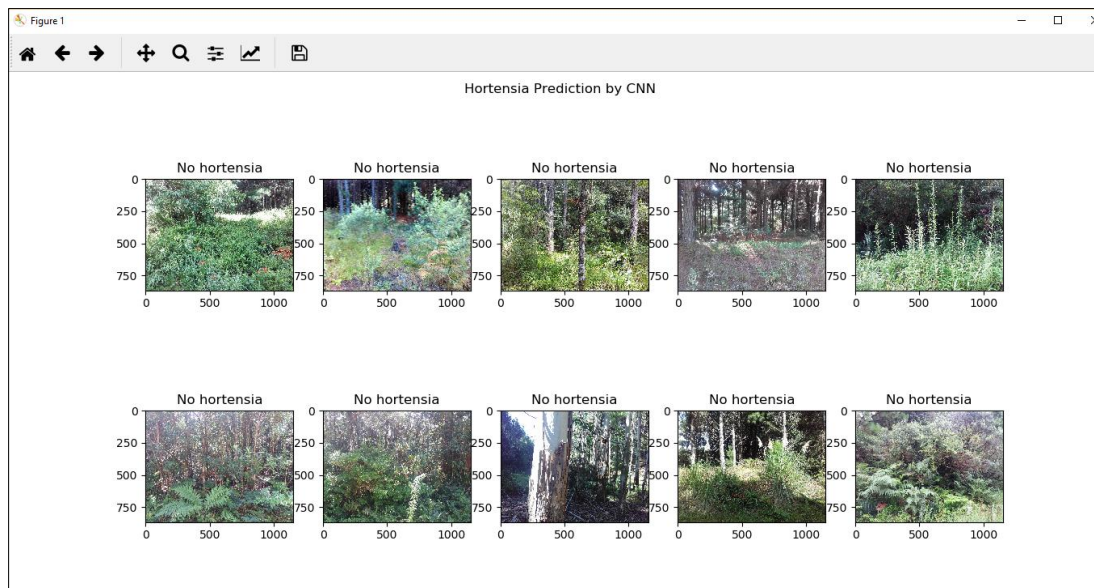
4.2 Testing the Trained CNNs

Once the CNNs have been trained to a high validation accuracy, without much over-fitting, the next step is to see whether the network is able to predict accurate results for data it hasn't previously encountered. For this purpose, 20 images had been saved in the test dataset of Hortensia, and 16 images had been saved in the test dataset of Pneumonia X-rays. At no point in its training have the CNN's ever come across these images.

The predictions of the Hortensia dataset trained CNN when tested on test images without the Hortensia flowers is given in Fig 4.5(a). The predictions by the CNN for test images with the Hortensia flowers is given in 4.5(b) and the accuracy for test images is seen in Fig 4.6. Similarly, the predictions of the Pneumonia X-Ray dataset trained CNN when tested on Normal X-Rays is given in Fig 4.7(a). The predictions by the CNN for Pneumonia X-Rays is given in 4.7(b) and the accuracy for test images is seen in Fig 4.8. This is performed by making use of predict and evaluate functions of the Keras library in Python.



(a)

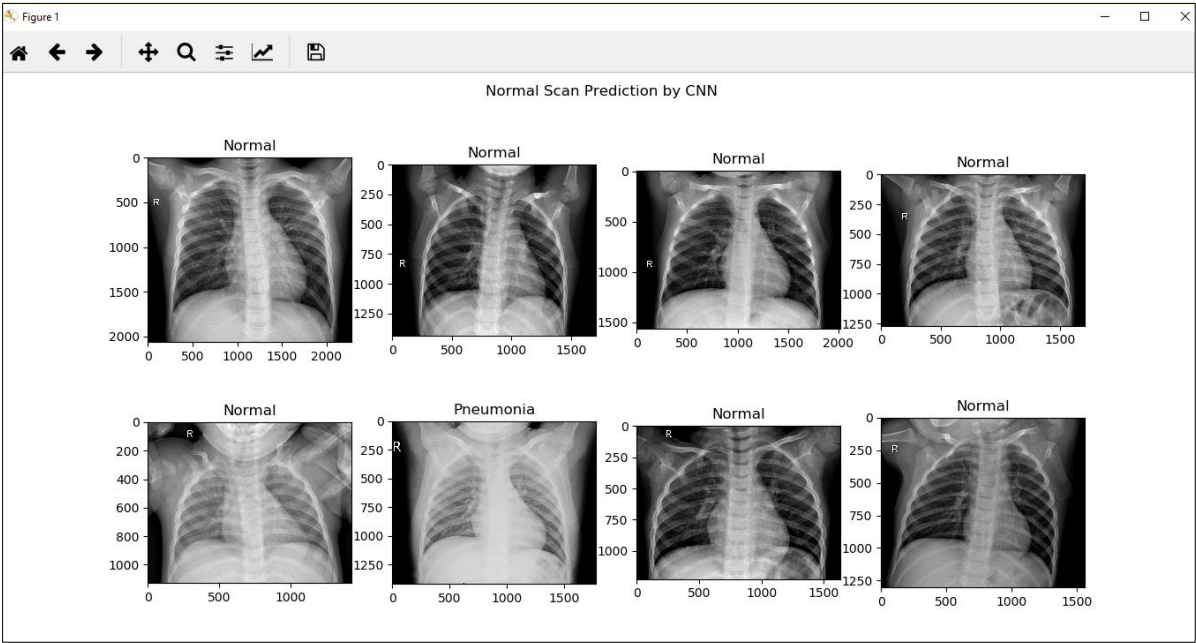


(b)

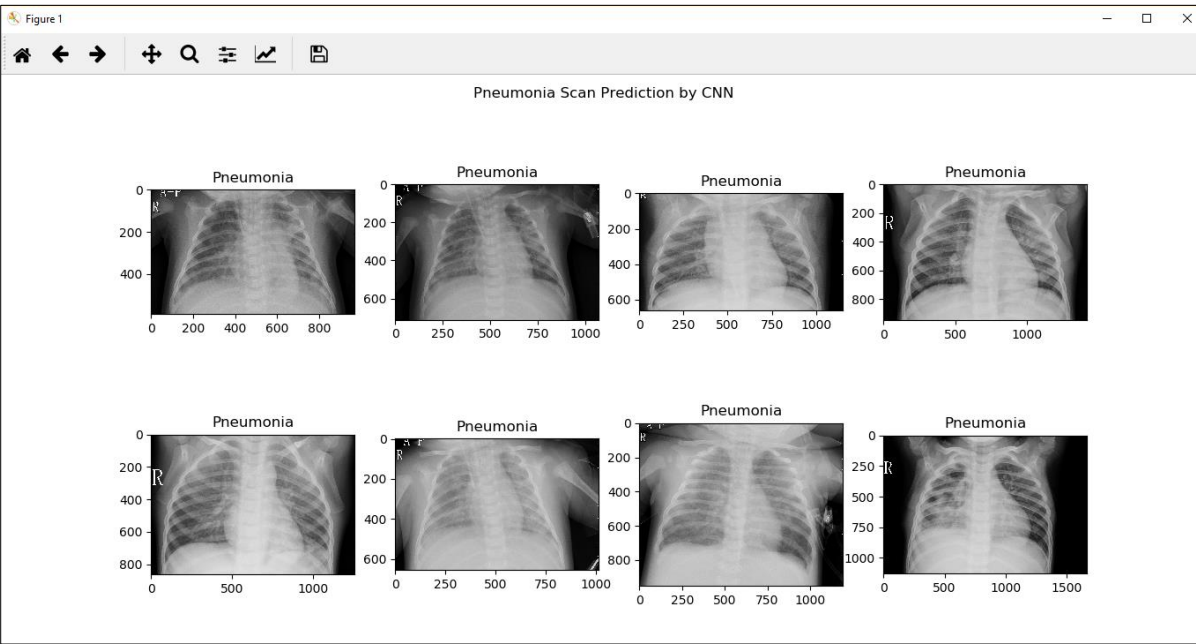
Fig 4.5 Predictions for (a) With flowers (b) Without flowers

```
Found 20 images belonging to 2 classes.  
the evaluation accuracy is 0.94
```

Fig 4.6 Testing accuracy Hortensia CNN



(a)



(b)

Fig 4.7 Predictions for (a) Normal X-Rays (b) Pneumonia X-Rays


```
Found 16 images belonging to 2 classes.  
the evaluation accuracy is 0.94
```

Fig 4.8 Testing accuracy Pneumonia X-Ray CNN

4.3 Deviation from Expected Results

In both the Hortensia test dataset and the Pneumonia test dataset, there was one image each that was incorrectly predicted. Fig 4.9 gives an enlarged view of the Hortensia image that the network incorrectly classified as No Hortensia. When compared with the other test images, it is evident that the Hortensia flowers are significantly smaller in size and more hidden. All images are resized into smaller images before being processed by the CNN as using pictures in their original sizes will involve far too much computation. As a result, when already small Hortensia flowers are resized into a smaller image, they will become insignificant to the network leading to an erroneous classification.

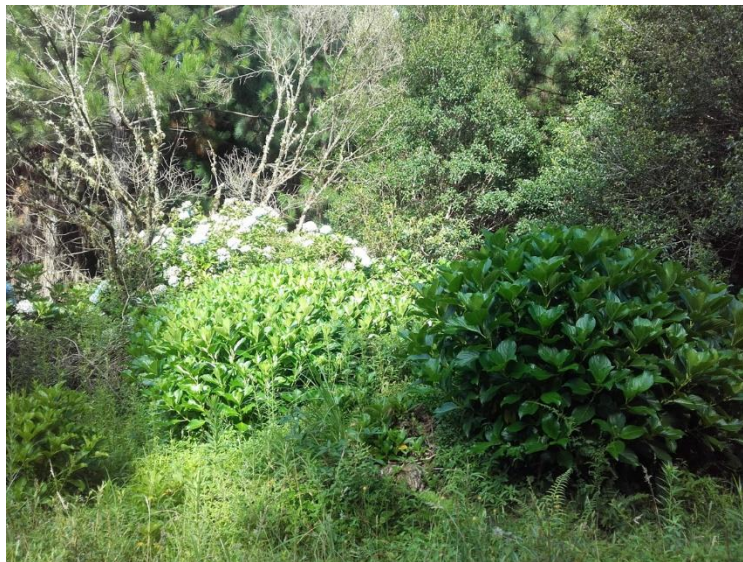


Fig 4.9 Incorrectly classified Hortensia test image

CNNs trained on medical images are slightly more prone to misclassification. This is because diseases do not look alike on all human bodies. If we have a look at the incorrectly classified image in the X-Ray dataset, shown in Fig. 4.10 which gives a false positive for Pneumonia, it can be easily observed that the opacity of the lungs is much less than the other normal test X-Rays. This could be due to the nature of this patient's lungs. This level of error is expected, considering even a human being would have difficulty correctly predicting such an X-Ray. Too high an accuracy for CNN's trained on medical data will indicate over-fitting to the data.



Fig 4.10 Incorrectly classified X-Ray

This chapter looked at how the validation and training accuracy varied with the number of epochs while training 2 different CNNs with external datasets. It also looked at the way the validation and training loss decreased with training. It saw how well the trained networks were able to predict the classes for images they hadn't encountered previously, and what caused erroneous results. The X-Ray CNN was able to predict images to a very high degree of accuracy and can be deployed in Samsung Digital Radiography machines, based on its ability to predict the presence of Pneumonia in X-Rays.

CHAPTER 5

CONCLUSION AND FUTURE SCOPE OF WORK

This project set out to develop a program that could identify whether a lung X-Ray showed the presence of Pneumonia or not, so that it could be incorporated into Samsung's digital radiography machines. The intention was to get a diagnosis instantly after taking a scan that could be used to aid a physician in making the final diagnosis. This project is the beginning of Samsung integrating artificial intelligence into its medical diagnostic devices.

The project was built using convolutional neural networks. Initially a CNN with a simple architecture was trained on one of the elementary datasets provided within the Keras library of Python itself, called MNIST. After getting an idea of the way a CNN functions, another CNN was trained to detect the presence of Hortensia flowers in picture of nature. This involved an external dataset. Tuning of hyper-parameters, and changing the architecture was understood. Modifications were then made to this CNN to train it using the Pneumonia X-Ray Dataset which required a larger architecture with more trainable parameters to accommodate the subtle differences between the two classes.

5.1 Conclusion

This project developed a program that was able to distinguish between Pneumonia X-Rays and Normal X-Rays to a high degree of validation accuracy, of 90%, which surpasses existing models. Medical images are particularly difficult to obtain a good accuracy with, owing to a high degree of intra-class variation (variation within the class) and a low degree of inter- class variation (variation between the classes to be distinguished). Thus, the obtained accuracy comes close to rivalling physicians diagnosing the disease. While the dataset was still significantly small to allow the program to serve as a standalone diagnostic tool, its high accuracy level will allow it to serve as an excellent aid to physicians. With an increase in the size and variety of the dataset, its ability to work as a reliable diagnostic tool in itself will increase.

The high validation accuracy the network gained by learning from 5216 X-Rays is a good indicator of the power of the network. It has the ability to diagnose Pneumonia with nearly the same accuracy as that of a physician who has seen the same number of Pneumonia X-Rays in her lifetime. This shows how fast artificial intelligence algorithms are catching up with physicians when it comes to their diagnostic ability. As the size and variety of available datasets keeps increasing, powerful convolutional networks will soon be able to surpass a trained physician's diagnostic ability, leading to fewer cases of misdiagnosis and the problems that follow it.

5.2 Future Scope of Work

Incorporating artificial intelligence into medical diagnostic tools is a very new field which is constantly expanding. Samsung Research Institute, Noida has recently decided to venture in big data in healthcare. Pneumonia is one of the most common diseases present in India and limiting its misdiagnosis could greatly benefit the population. This project makes use of a dataset that has lung X-Rays from children between the ages of 1-5. Tying up with local hospitals to get Pneumonia scans for research purposes, to generate datasets of more age variation will lead to a larger reach for the project. The project in its current state might have difficulty in diagnosing Pneumonia in neonatal or adult patients owing to the lack of such images while training. Also, the larger the dataset is, the lesser the network is prone to over-fitting, leading to a very robust convolutional network. Thus, increasing the size and variety of the available Pneumonia datasets will lead to more powerful networks.

Different X-Ray machines capture different types of images. There may be variations in image shape, size and even colour. Samsung can collect data from all its digital radiography devices deployed across the world, with patient consent, which will all be in the same format. This can be used to train a network that doesn't train itself incorrectly owing to the difference in the nature of the data.

Further, this project can be expanded to encompass other lung diseases as well. Instead of it being a binary classification problem, the network can be trained for multi-class classification of other diseases the respiratory tract infection, lung cancer, bronchitis etc. if such datasets become available in the future, with a very minor change in the structure. This can be done by changing the final activation to a soft-max function to give the probability of each class.

REFERENCES

- [1]. Chen, Hsinchun, Roger HL Chiang "Business intelligence and analytics: from big data to big impact." *MIS quarterly* (2012): 1165-1188.
- [2]. Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.
- [3]. John Walker, Saint. "Big data: A revolution that will transform how we live, work, and think." (2014): 181-183
- [4]. Facebook Convolutional Neural Networks, <https://code.facebook.com/>
- [5]. Apple Machine learning, <https://machinelearning.apple.com/>
- [6]. SONY Big Data Applications, <http://bigdata-madesimple.com>

ANNEXURES

Code

1. Code for X-Ray Dataset trained CNN

Xray_start.py – Code to load the dataset, check its size and view samples

```
from os import environ, chdir
import scipy.ndimage as sp
import os
import numpy as np
import random
import matplotlib.pyplot as plt

from keras.preprocessing.image import ImageDataGenerator
from keras import models, layers, optimizers, callbacks

environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
chdir(r'C:\skr\xray')
train_images_normal = os.listdir(r'data1\train\NORMAL')
train_images_xray = os.listdir(r'data1\train\PNEUMONIA')
test_images_normal = os.listdir(r'data1\test\NORMAL')
test_images_xray = os.listdir(r'data1\test\PNEUMONIA')
val_images_normal = os.listdir(r'data1\val\NORMAL')
val_images_xray = os.listdir(r'data1\val\PNEUMONIA')
print('X-RAY data-set size')
print('No. of images in Train data-set=',
      len(train_images_xray)+len(train_images_normal))
print('No. of images in Test data-set=',
      len(test_images_xray)+len(test_images_normal))
print('No. of images in Valid data-set=',
      len(val_images_xray)+len(val_images_normal))

print('Normal Lung Samples')
sample_normal = random.sample(train_images_normal,6)
f,ax = plt.subplots(2,3,Figsize=(15,9))
for i in range(0, 6):
    im = sp.imread(r'data1/train/NORMAL/' + sample_normal[i])
    print(im.shape)
    ax[i // 3, i % 3].imshow(im, cmap=plt.cm.gray)
    ax[i // 3, i % 3].axis('on')
f.suptitle('Normal Lungs')
plt.show()
print('Pneumonia Lung Samples')
sample_normal1 = random.sample(train_images_xray,6)
f,ax = plt.subplots(2,3,Figsize=(15,9))
for i in range(0, 6):
    im = sp.imread(r'data1/train/PNEUMONIA/' + sample_normal1[i])
    print(im.shape)
    ax[i // 3, i % 3].imshow(im, cmap=plt.cm.gray)
    ax[i // 3, i % 3].axis('on')
f.suptitle('Pneumonia Lungs')
plt.show()
```

Xray_train.py– Code to design the network, select loss and optimiser, set call-backs, train the network, save the model and plot accuracy and loss graphs.

```
from os import environ, chdir
from keras.preprocessing.image import ImageDataGenerator
from keras import models, layers, optimizers, callbacks
from keras.models import Sequential
from keras.layers import Conv2D,Dense,Flatten,Dropout,MaxPooling2D
import matplotlib.pyplot as plt
```

```

from skimage import img_as_ubyte

environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
chdir(r'C:\skr\xray')

# Setting Image Generators

train_idg = ImageDataGenerator(rescale=1./255, shear_range=0.2, zoom_range=0.2,
horizontal_flip=True)
train_g = train_idg.flow_from_directory(directory=r'data1\train', target_size=(150,
150),
                                     class_mode='binary', batch_size=32)
valid_idg = ImageDataGenerator(rescale=1./255)
valid_g = valid_idg.flow_from_directory(directory=r'data1\val', target_size=(150,
150),
                                     class_mode='binary', batch_size=32)

# CNN Architecture

my_model = Sequential()
my_model.add(Conv2D(64, (3,3), input_shape=(150,150,3), activation='relu'))
my_model.add(Conv2D(64, (3,3), activation='relu'))
my_model.add(MaxPooling2D(pool_size=(2,2)))
my_model.add(Dropout(0.2))
my_model.add(Conv2D(32, (3,3), activation='relu'))
my_model.add(Conv2D(32, (3,3), activation='relu'))
my_model.add(MaxPooling2D(pool_size=(2,2)))
my_model.add(Dropout(0.2))
my_model.add(Conv2D(16, (3,3), activation='relu'))
my_model.add(Conv2D(16, (3,3), activation='relu'))
my_model.add(MaxPooling2D(pool_size=(2,2)))
my_model.add(Dropout(0.2))
my_model.add(Flatten())
my_model.add(Dense(units=128, activation='relu'))
my_model.add(Dropout(0.2))
my_model.add(Dense(units=1, activation='sigmoid'))

print(my_model.summary())

# Model Loss Function and Optimiser

my_model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

# Callbacks

check_p = callbacks.ModelCheckpoint(filepath='xray9_cnn_{val_acc:.2f}.h5',
monitor='val_acc',
                                verbose=1, save_best_only=True,
save_weights_only=False)
reduce_lr = callbacks.ReduceLROnPlateau(monitor='loss', factor=0.1, patience=2,
cooldown=2, min_lr=0.00001, verbose=1)
call_l = [check_p, reduce_lr]

# Training Options

fit = my_model.fit_generator(generator=train_g, steps_per_epoch=163, epochs=13,
verbose=1, callbacks=call_l,
                           validation_data=valid_g, validation_steps=19)

print(fit.history.keys())
# summarize history for accuracy
plt.plot(fit.history['acc'])
plt.plot(fit.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper left')

```

```

plt.show()

# summarize history for loss
plt.plot(fit.history['loss'])
plt.plot(fit.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper left')
plt.show()

# Saving Model
my_model.save(filepath=r'xray9_cnn.h5', overwrite=True)

```

Xray_test.py – Code to load the trained model, display predictions for test images and calculate the test accuracy

```

import os
from keras.models import load_model
from keras.preprocessing.image import ImageDataGenerator
from numpy import set_printoptions
import random
import matplotlib.pyplot as plt
import scipy.ndimage as sp

from keras.preprocessing import image

set_printoptions(precision=4, suppress=True)
# Initial Settings

os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
os.chdir(r'C:\skr\xray')

# Loading Model

my_model = load_model(filepath=r'xray6_cnn_0.91.h5')

print(my_model.summary())

# Weights and biases

'''
print('Hydrangea last layer weights')
print(my_model.get_weights()[-2])
'''

# Evaluation

eval_idg = ImageDataGenerator(rescale=1./255)
eval_g = eval_idg.flow_from_directory(directory=r'data1\test', target_size=(150,
150),
                                   class_mode='binary', batch_size=16,
shuffle=False)
(eval_loss, eval_acc) = my_model.evaluate_generator(generator=eval_g, steps=1)
print('the evaluation accuracy is {:.42f}'.format(eval_acc))

# Individual Prediction
'''
pre_idg = eval_idg
pre_g = eval_g

pre = my_model.predict_generator(generator=pre_g, steps=1)
print(pre_g.filenames, '\n')
print(pre_g.class_indices, '\n')
print(pre[0:8]<0.5, '\n')
'''

```

```

print(pre[8:16]>0.5, '\n')
'''

test_images_normal = os.listdir(r'data1\test\NORMAL')
test_images_xray = os.listdir(r'data1\test\PNEUMONIA')

f, ax = plt.subplots(2, 4, Figsize=(15, 9))
for i in range(0, 8):

    name = test_images_normal[i]
    im = sp.imread(r'C:/skr/xray/data1/test/NORMAL/' + name)
    ax[i // 4, i % 4].imshow(im, cmap=plt.cm.gray)
    ax[i // 4, i % 4].axis('on')
    x = image.load_img(r'C:/skr/xray/data1/test/NORMAL/' + name, target_size=(150,
150))
    x = image.img_to_array(x)
    x = x / 255
    pred = my_model.predict(x.reshape((1, 150, 150, 3)))
    if pred < 0.5:
        ax[i // 4, i % 4].set_title('Normal')
    else:
        ax[i // 4, i % 4].set_title('Pneumonia')

f.suptitle('Normal Scan Prediction by CNN')
plt.show()

f, ax = plt.subplots(2, 4, Figsize=(15, 9))
for i in range(0, 8):

    name = test_images_xray[i]
    im = sp.imread(r'C:/skr/xray/data1/test/PNEUMONIA/' + name)
    ax[i // 4, i % 4].imshow(im, cmap=plt.cm.gray)
    ax[i // 4, i % 4].axis('on')
    x = image.load_img(r'C:/skr/xray/data1/test/PNEUMONIA/' + name,
target_size=(150, 150))
    x = image.img_to_array(x)
    x = x / 255
    pred = my_model.predict(x.reshape((1, 150, 150, 3)))
    if pred < 0.5:
        ax[i // 4, i % 4].set_title('Normal')
    else:
        ax[i // 4, i % 4].set_title('Pneumonia')

f.suptitle('Pneumonia Scan Prediction by CNN')
plt.show()

```

2. Code for Hortensia Dataset trained CNN

Hydrangea_start.py- Code to load the dataset, check its size and view samples

```

from os import environ, chdir
import scipy.ndimage as sp
import os
import numpy as np
import random
import matplotlib.pyplot as plt
import random

environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
chdir(r'C:\skr\hydrangea')
train_images_hortensia = os.listdir(r'data2\train\hortensia')
train_images_nature=os.listdir(r'data2\train\nature')
test_images_hortensia = os.listdir(r'data2\eval\hortensia')
test_images_nature=os.listdir(r'data2\eval\nature')
valid_images_hortensia = os.listdir(r'data2\valid\hortensia')
valid_images_nature=os.listdir(r'data2\valid\nature')

```

```

print('Size of Hortensia Dataset')
print('No. of images in Train data-set=',
len(train_images_hortensia)+len(train_images_nature))
print('No. of images in Test data-set=',
len(test_images_hortensia)+len(test_images_nature))
print('No. of images in Valid data-set=',
len(valid_images_hortensia)+len(valid_images_nature))

print('Hortensia Samples')
sample_normal = random.sample(train_images_hortensia,6)
f,ax = plt.subplots(2,3,Figsize=(15,9))
for i in range(0, 6):
    im = sp.imread(r'data2/train/hortensia/' + sample_normal[i])
    print(im.shape)
    ax[i // 3, i % 3].imshow(im)
    ax[i // 3, i % 3].axis('on')
f.suptitle('Hortensia')
plt.show()
print('No Hortensia Samples')
sample_normal1 = random.sample(train_images_nature,6)
f,ax = plt.subplots(2,3,Figsize=(15,9))
for i in range(0, 6):
    im = sp.imread(r'data2/train/nature/' + sample_normal1[i])
    print(im.shape)
    ax[i // 3, i % 3].imshow(im)
    ax[i // 3, i % 3].axis('on')

f.suptitle('No Hortensia')
plt.show()

```

hydrangea_train.py - Code to design the network, select loss and optimiser, set call-backs, train the network, save the model and plot accuracy and loss graphs

```

from os import environ, chdir
from keras.preprocessing.image import ImageDataGenerator
from keras import models, layers, optimizers, callbacks
import matplotlib.pyplot as plt

environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
chdir(r'C:\skr\hydrangea')

# Setting Image Generators

train_idg = ImageDataGenerator(rescale=1./255, zoom_range=[1.0, 1.25],
width_shift_range=0.1,
                                height_shift_range=0.1, fill_mode='reflect')
train_g = train_idg.flow_from_directory(directory=r'data2\train', target_size=(100,
100),
                                class_mode='binary', batch_size=125,
shuffle=True)
valid_idg = ImageDataGenerator(rescale=1./255)
valid_g = valid_idg.flow_from_directory(directory=r'data2\valid', target_size=(100,
100),
                                class_mode='binary', batch_size=125,
shuffle=True)

# CNN Architecture

my_model = models.Sequential()
my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1),
input_shape=(100, 100, 3)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1),
input_shape=(100, 100, 3)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.MaxPooling2D(pool_size=(2, 2)))
my_model.add(layers.Dropout(rate=0.3))

```



```

my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.MaxPooling2D(pool_size=(2, 2)))
my_model.add(layers.Dropout(rate=0.3))
my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.MaxPooling2D(pool_size=(2, 2)))
my_model.add(layers.Conv2D(filters=16, kernel_size=(2, 2), strides=(1, 1)))
my_model.add(layers.Activation('relu'))
my_model.add(layers.MaxPooling2D(pool_size=(2, 2)))
my_model.add(layers.Flatten())
my_model.add(layers.Dropout(rate=0.3))
my_model.add(layers.Dense(units=10))
my_model.add(layers.Activation('relu'))
my_model.add(layers.Dense(units=1))
my_model.add(layers.Activation('sigmoid'))

print(my_model.summary())

# Model Loss Function and Optimiser

compile1 = my_model.compile(optimizer=optimizers.sgd(lr=0.15),
loss='binary_crossentropy', metrics=['accuracy'])

# Callbacks

check_p = callbacks.ModelCheckpoint(filepath='hydrangeal_cnn_{val_acc:.2f}.h5',
monitor='val_acc',
                                verbose=1, save_best_only=True,
save_weights_only=False)
reduce_lr = callbacks.ReduceLROnPlateau(monitor='val_acc', patience=5, factor=0.95,
verbose=1, cooldown=2)
call_l = [check_p, reduce_lr]

# Training Options

fit = my_model.fit_generator(generator=train_g, steps_per_epoch=22, epochs=100,
verbose=1, callbacks=call_l,
                        validation_data=valid_g, validation_steps=4)

print(fit.history.keys())
# summarize history for accuracy
plt.plot(fit.history['acc'])
plt.plot(fit.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper left')
plt.show()

# summarize history for loss
plt.plot(fit.history['loss'])
plt.plot(fit.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'valid'], loc='upper left')
plt.show()
# Saving Model

my_model.save(filepath=r'hydrangeal_cnn.h5', overwrite=True)

```

hydrangea_test.py - Code to load the trained model, display predictions for test images and calculate the test accuracy

```

import os
from keras.models import load_model
from keras.preprocessing.image import ImageDataGenerator
from numpy import set_printoptions
import matplotlib.pyplot as plt
import random
import scipy.ndimage as sp
from skimage import img_as_ubyte

from keras.preprocessing import image
set_printoptions(precision=4, suppress= True)
# Initial Settings

os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
os.chdir(r'C:\skr\hydrangea')

# Loading Model

my_model = load_model(filepath=r'hydrangea_cnn_0.92.h5')
'''
print(my_model.summary())
'''

test_images_hortensia = os.listdir(r'data2\eval\hortensia')
test_images_nature = os.listdir(r'data2\eval\nature')

eval_idg = ImageDataGenerator(rescale=1./255)
eval_g = eval_idg.flow_from_directory(directory=r'data2\eval', target_size=(100,
100),
                                class_mode='binary', batch_size=16,
shuffle=False)
(eval_loss, eval_acc) = my_model.evaluate_generator(generator=eval_g, steps=1)
print('the evaluation accuracy is {:.2f}'.format(eval_acc))

f, ax = plt.subplots(2, 5, figsize=(15, 9))
for i in range(0, 10):

    name = test_images_hortensia[i]
    im=sp.imread(r'C:/skr/hydrangea/data2/eval/hortensia/' + name)
    ax[i//5, i % 5].imshow(im, cmap=plt.cm.gray)
    ax[i//5, i % 5].axis('on')
    x = image.load_img(r'C:/skr/hydrangea/data2/eval/hortensia/' + name,
target_size=(100, 100))
    x = image.img_to_array(x)
    x = x / 255
    pred = my_model.predict(x.reshape((1, 100, 100, 3)))
    if pred < 0.5:
        ax[i//5, i % 5].set_title('hortensia')
    else:
        ax[i//5, i % 5].set_title('No hortensia')

f.suptitle('Hortensia Prediction by CNN')
plt.show()

f, ax = plt.subplots(2, 5, figsize=(15, 9))
for i in range(0, 10):

    name = test_images_nature[i]
    im = sp.imread(r'C:/skr/hydrangea/data2/eval/nature/' + name)
    ax[i//5, i % 5].imshow(im, cmap=plt.cm.gray)
    ax[i//5, i % 5].axis('on')
    x = image.load_img(r'C:/skr/hydrangea/data2/eval/nature/' + name,
target_size=(100, 100))
    x = image.img_to_array(x)
    x = x / 255
    pred = my_model.predict(x.reshape((1, 100, 100, 3)))

```

```
if pred < 0.5:
    ax[i//5, i % 5].set_title('hortensia')
else:
    ax[i//5, i % 5].set_title('No hortensia')

f.suptitle('Hortensia Prediction by CNN')
plt.show()
```

PROJECT DETAILS

<i>Student Details</i>			
Student Name	Shreya K R		
Register Number	140907759	Section / Roll No	D-60
Email Address	shreyakr96@gmail.com	Phone No (M)	9663598199
<i>Project Details</i>			
Project Title	Convolutional Neural Networks based Pneumonia Diagnosis		
Project Duration	5 months	Date of reporting	8-01-2018
Expected date of completion of project	26-05-2018		
<i>Organization Details</i>			
Organization Name	Samsung Research Institute, Noida		
Full postal address with pin code	Samsung Research Institute, Plot No. C 28-29, Tower D, Logix Cyber Park, Overseas Lane, C Block, Sector 62, Noida, Uttar Pradesh 201301		
Website address	http://www.samsung.com		
<i>Supervisor Details</i>			
Supervisor Name	Mr. Vijay Kumar Gupta		
Designation	Deputy General Manager		
Full contact address with pin code	Samsung Research Institute, Plot No. C 28-29, Tower D, Logix Cyber Park, Overseas Lane, C Block, Sector 62, Noida, Uttar Pradesh 201301		
Email address	g.vijay@samsung.com	Phone No (M)	
<i>Internal Guide Details</i>			
Faculty Name	Ms. Navya KT		
Full contact address with pin code	Dept. of E&C Engg., Manipal Institute of Technology, Manipal – 576 104 (Karnataka State), INDIA		
Email address	navya.kt@manipal.edu		